

# netkit lab

## Traffic Engineering with MPLS for Linux

<b>Version</b>	1.1
<b>Author(s)</b>	F. Di Ciccio, F. Antonini (Kasko Networks S.r.l.) Reviewed by M. Rimondini (Roma Tre University)
<b>E-mail</b>	fra.dix87@gmail.com, fabio.antonini@kaskonetworks.it
<b>Web</b>	<a href="http://www.kaskonetworks.it/">http://www.kaskonetworks.it/</a>
<b>Description</b>	An example with 5 routers (2 lers, 3 lsrs) and 2 hosts to show link protection on a link and node protection on a router, with MPLS

# copyright notice

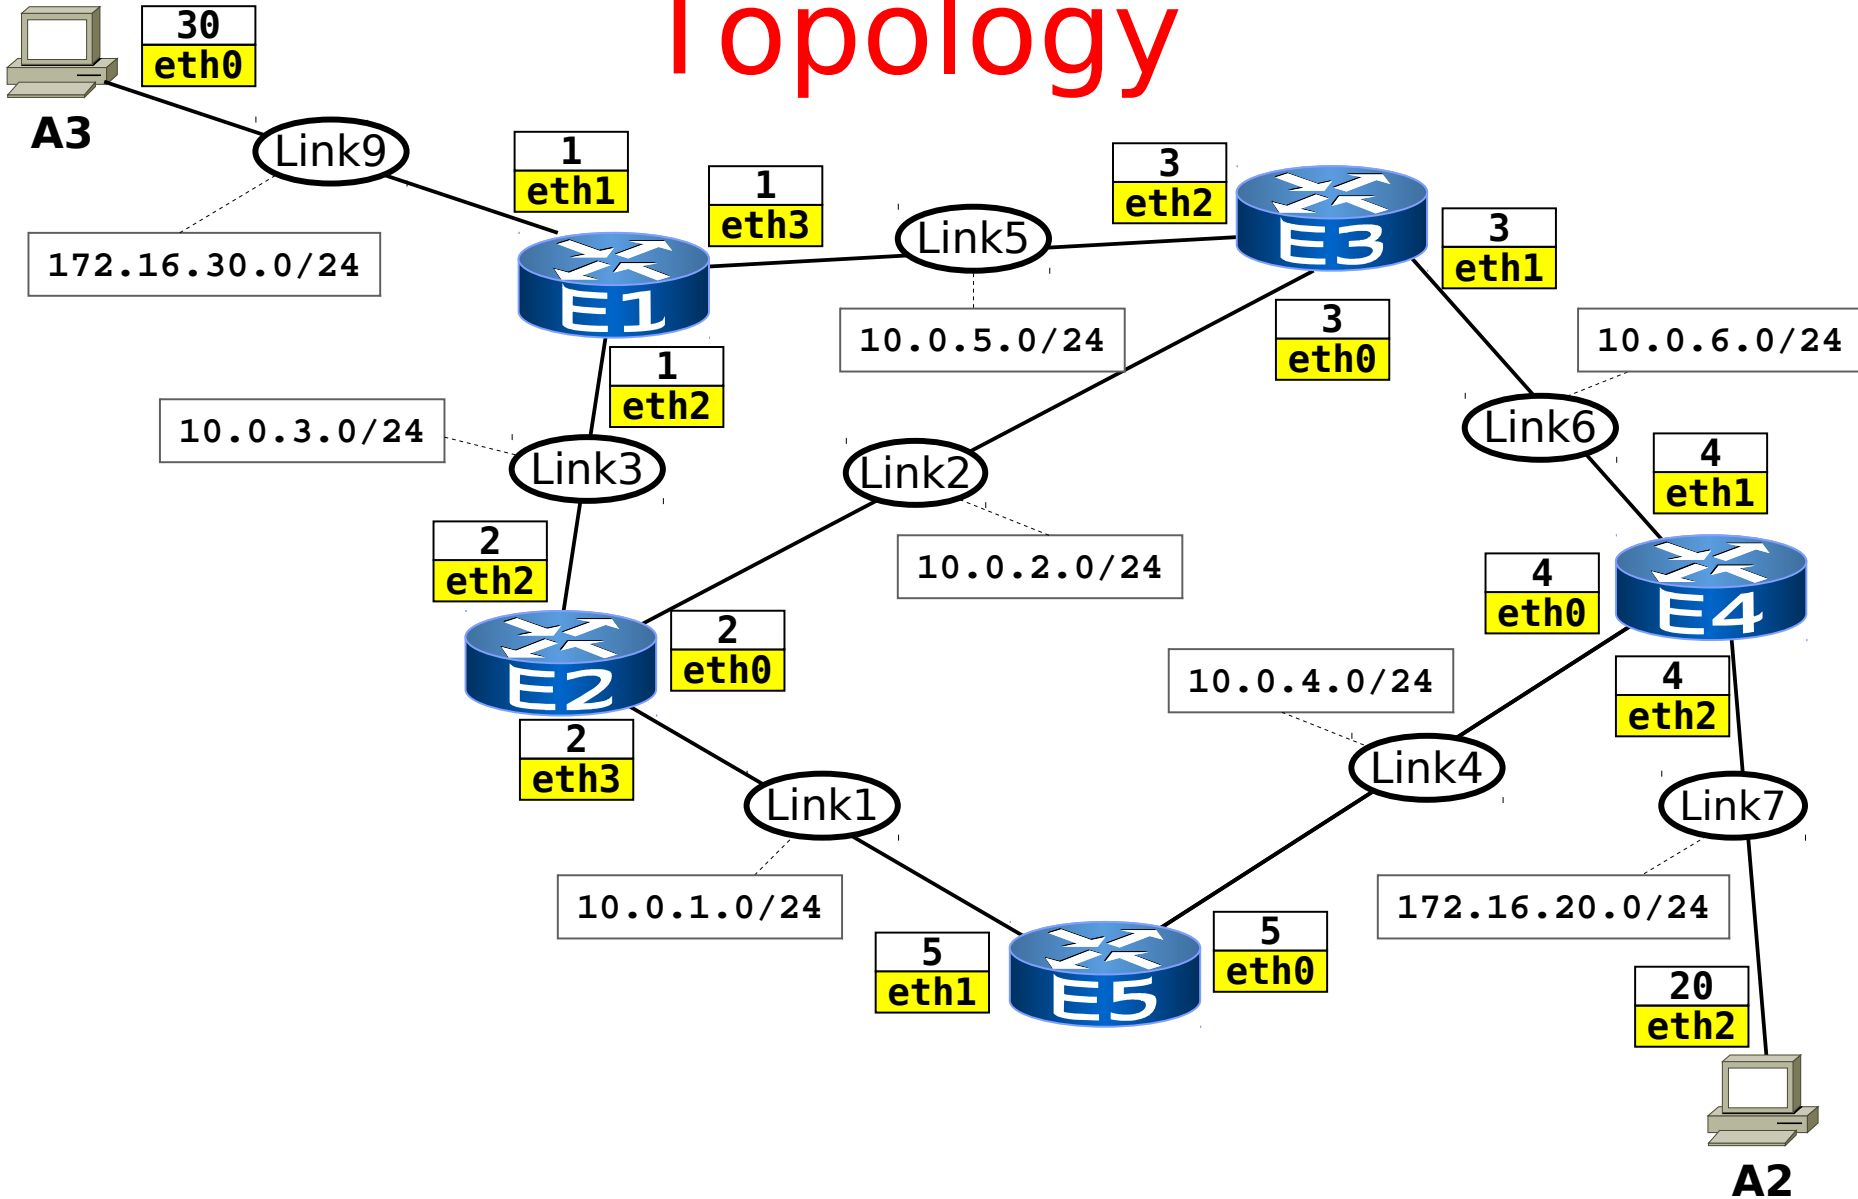
- All the pages/slides in this presentation, including but not limited to, images, photos, animations, videos, sounds, music, and text (hereby referred to as “material”) are protected by copyright.
- This material, with the exception of some multimedia elements licensed by other organizations, is property of the authors and/or organizations appearing in the first slide.
- This material, or its parts, can be reproduced and used for didactical purposes within universities and schools, provided that this happens for non-profit purposes.
- Information contained in this material cannot be used within network design projects or other products of any kind.
- Any other use is prohibited, unless explicitly authorized by the authors on the basis of an explicit agreement.
- The authors assume no responsibility about this material and provide this material “as is”, with no implicit or explicit warranty about the correctness and completeness of its contents, which may be subject to changes.
- This copyright notice must always be redistributed together with the material, or its portions.

# The MPLS Lab

- This lab is (highly) inspired to the example “MPLS-Linux labs – 4.2 Node Protection” from [1]
- Goal: achieving fast & seamless rerouting in case of link/node failure using MPLS
- Approach: emulate fast rerouting by changing label swapping rules on the fly upon detecting link/node failures

[1] [http://ontwerpen1.khlim.be/~lrutten/cursussen/comm2/mpls-linux-docs/4-2-node\\_protection.html](http://ontwerpen1.khlim.be/~lrutten/cursussen/comm2/mpls-linux-docs/4-2-node_protection.html)


# Topology



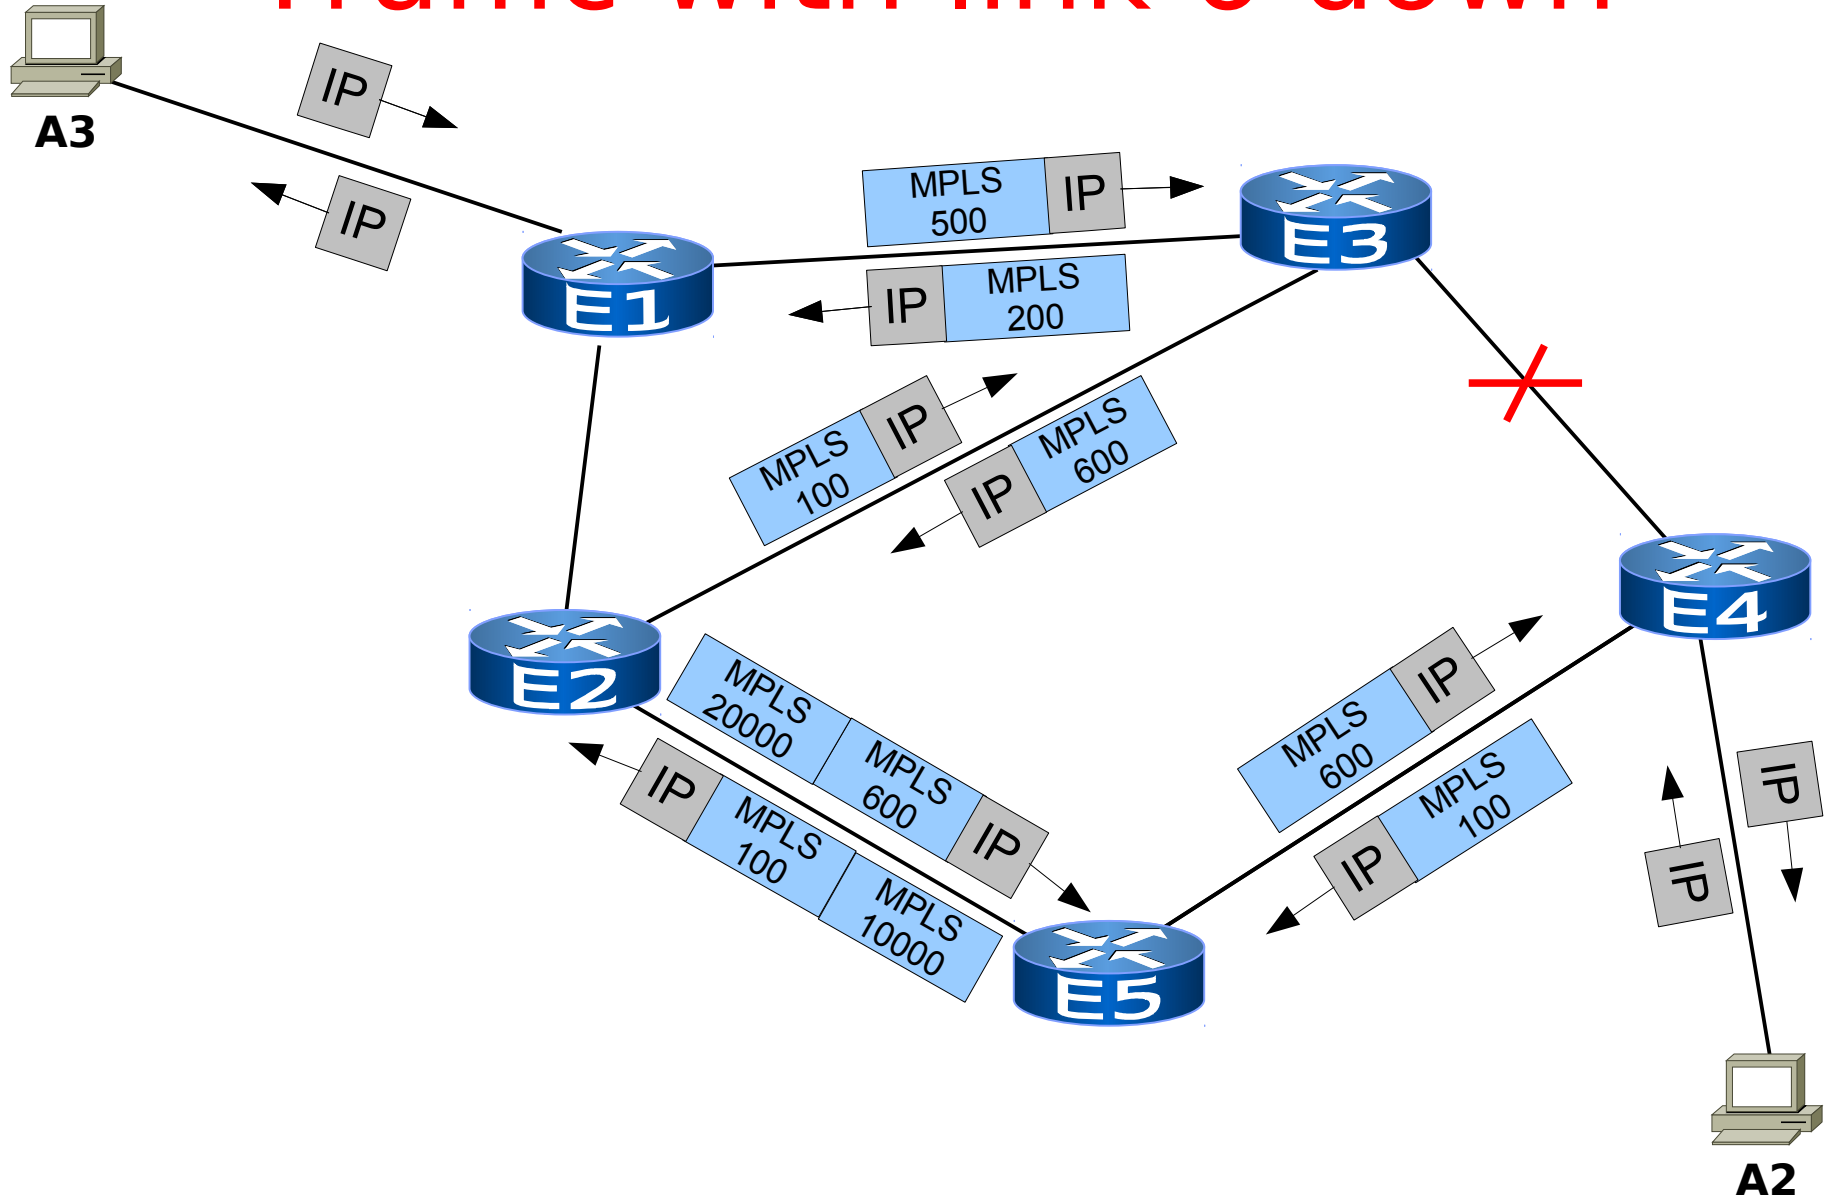
# Network topology

- A3 and A2 are connected by an MPLS network (E1, E2, E3, E4, E5)
  - Note: there are no VPNs, i.e., we use a single label
- For the traffic from A2 to A3 routers are configured to choose the “best” route (E4, E3, E1)
- If link 6 fails, all traffic between A2 and A3 is switched to a backup path (E4, E5, E2, E3, E1)
- If node E3 fails, all traffic between A2 and A3 is switched to a backup route (E4, E5, E2, E1)

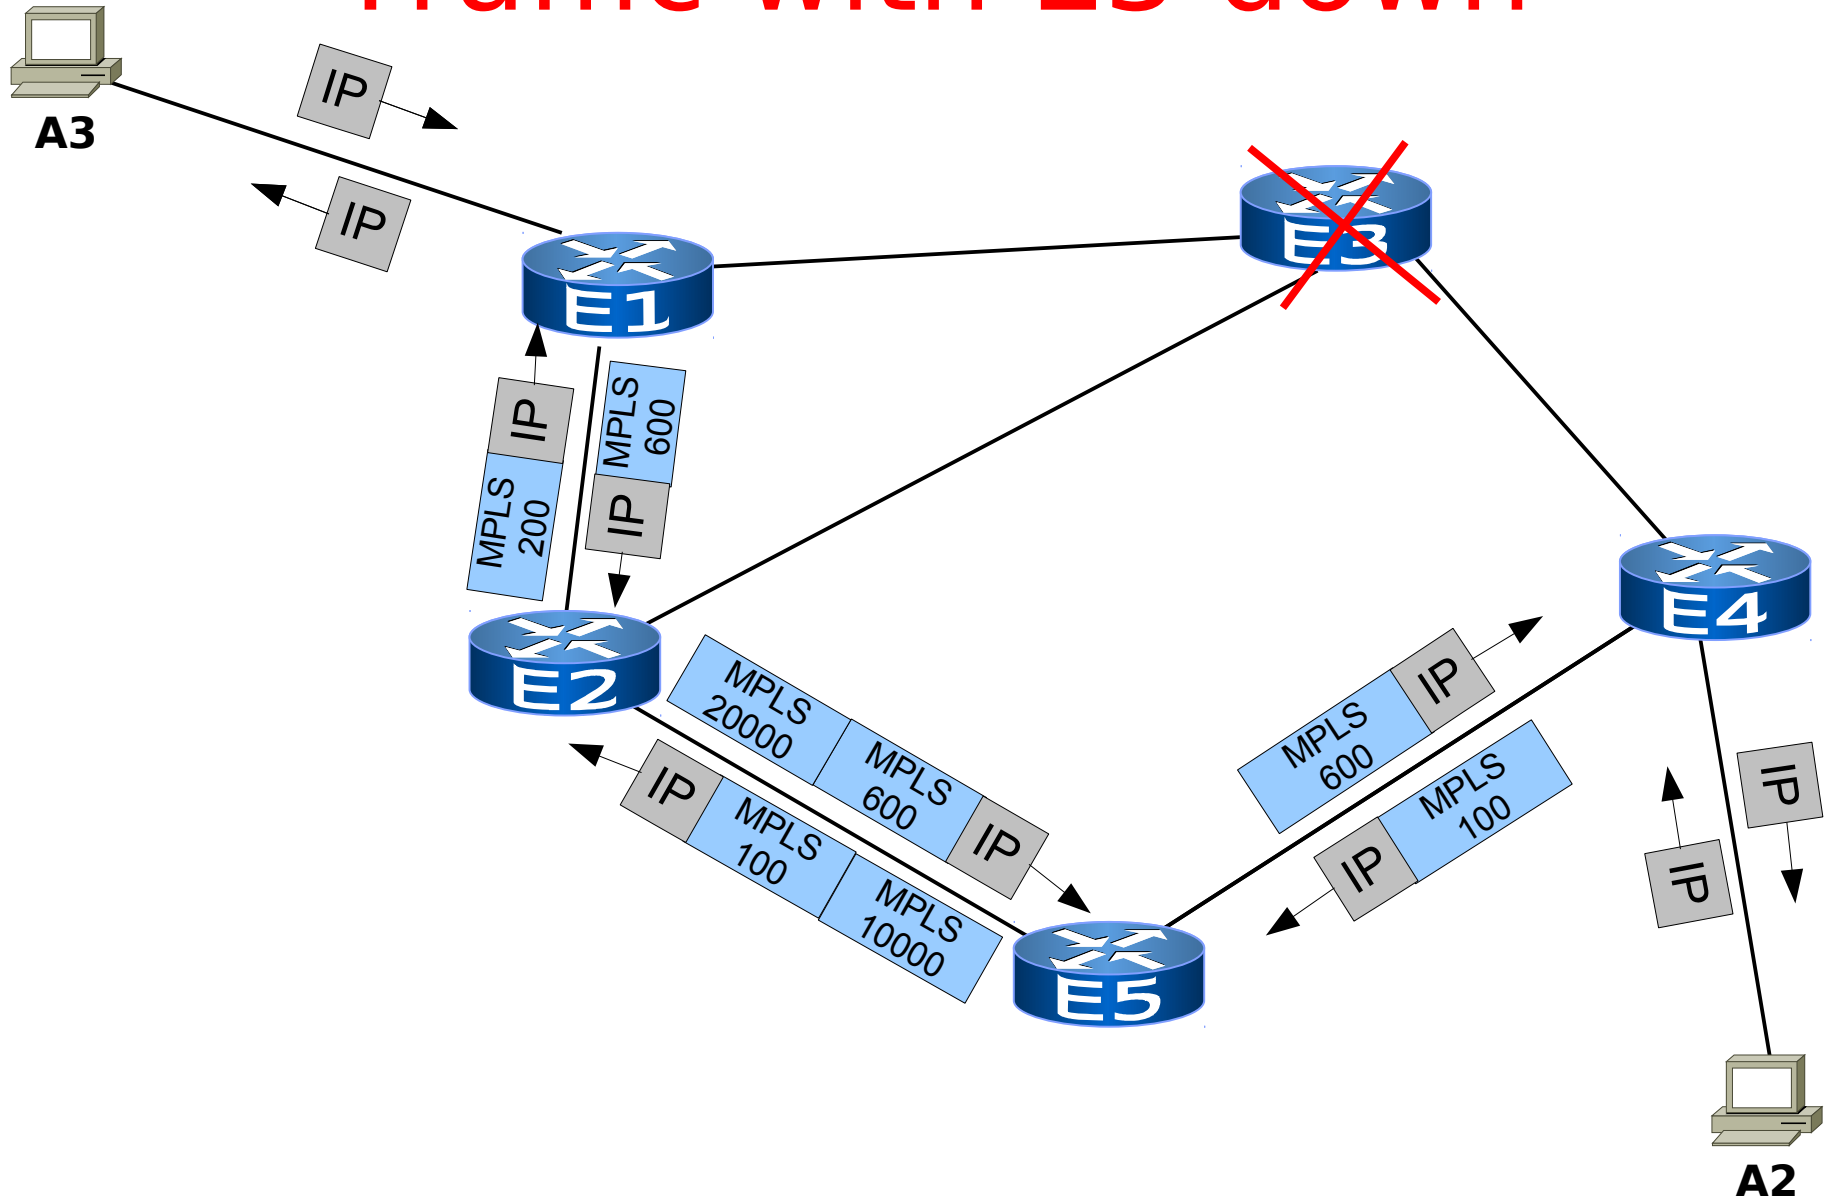
# Normal traffic



# Traffic with link 6 down



# Traffic with E3 down





# Basic configuration on end hosts

## A2.startup

```
ifconfig eth2 down
ifconfig eth2 172.16.20.20 netmask 255.255.255.0 up
ip route add 172.16.10.0/24 via 172.16.20.4 dev eth2
ip route add 172.16.30.0/24 via 172.16.20.4 dev eth2
```

## A3.startup

```
ifconfig eth0 down
ifconfig eth0 172.16.30.30 netmask 255.255.255.0 up
ip route add 172.16.20.0/24 dev eth0 via 172.16.30.1
ip route add 172.16.10.0/24 dev eth0 via 172.16.30.1
```

# Router settings

- All routers load MPLS kernel modules at startup
- MPLS configuration on E5 is in E5.startup as usual
- MPLS configurations for E1-E4 are in external scripts found in shared/root/ (therefore /root inside virtual machines)
  - These scripts take care of redirecting traffic as soon as failures are detected
  - They periodically ping to check which routes are still available
    - best: via E4, E3, E1
    - medium: via E4, E5, E2, E3, E1
    - low: via E4, E5, E2, E1

# MPLS configuration on Ier E1

## MPLS\_on\_E1.sh

#A2->A3 default route

```
mpls labelspace set dev eth3 labelspace 0
```

```
mpls ilm add label gen 200 labelspace 0
```

add an entry on ilm  
table in order to list  
(and pop) the incoming  
label (200)

Enable eth3  
to receive  
MPLS traffic

#A2->A3 backup route

```
mpls labelspace set dev eth2 labelspace 0
```

# MPLS configuration on Ier E1

## MPLS\_on\_E1.sh

#A3->A2 default route

```
var_best=`mpls nhlfe add key 0 instructions push gen 500 nexthop  
eth3 ipv4 10.0.5.3|grep key|cut -c 17-26`
```

```
ip route add 172.16.20.0/24 via 10.0.5.3 mpls $var_best
```

New nhlfe  
entry

sequential number identifying  
the entry (0="new entry": a  
number will be automatically  
assigned)

#A3->A2 backup route

```
var_low=`mpls nhlfe add key 0 instructions push gen 600 nexthop  
eth2 ipv4 10.0.3.2|grep key|cut -c 17-26`
```

# MPLS configuration on Ier E1

## MPLS\_on\_E1.sh

#A3->A2 default route

```
var_best=`mpls nhlf add key 0 instructions push gen 500 nexthop  
eth3 ipv4 10.0.5.3|grep key|cut -c 17-26`
```

```
ip route add 172.16.20.0/24 via 10.0.5.1 mpls $var_best
```

Push a label of  
type "gen" and  
value 500...

...and forward  
the packet to a  
certain router

#A3->A2 backup route

```
var_low=`mpls nhlf add key 0 instructions push gen 600 nexthop  
eth2 ipv4 10.0.3.2|grep key|cut -c 17-26`
```

# MPLS configuration on Ier E1

## MPLS\_on\_E1.sh

#A3->A2 default route

```
var_best=`mpls nhlfe add key 0 instructions push gen 500 nexthop  
eth3 ipv4 10.0.5.3|grep key|cut -c 17-26`
```

```
ip route add 172.16.20.0/24 via 10.0.5.3 mpls $var_best
```

label binding: instruct the router to use the previously created nhlfe to forward the packet

#A3->A2 backup route

```
var_low=`mpls nhlfe add key 0 instructions push gen 500 nexthop  
eth2 ipv4 10.0.3.2|grep key|cut -c 17-26`
```

this is the key returned by the previous mpls nhlfe add command  
note: here we are defining a fec

# MPLS configuration on Ier E1

Loop

MPLS\_on\_E1.sh

```
route="best"
while [ "1" = "1" ]
do
  case $route in
    best)
      status=`/usr/sbin/hping3 -c 1 --icmp 10.0.5.3 2>&1
              |grep "100% packet loss\|Host Unreachable"`
      if [ ! -z "$status" ] ;then
        # E3 is down switch to low route
        route="low";
        ip route del 172.16.20.0/24 via 10.0.5.3 mpls $var_best
        ip route add 172.16.20.0/24 via 10.0.3.2 mpls $var_low
      fi
    ;;
    low)
      status=`/usr/sbin/hping3 -c 1 --icmp 10.0.5.3 2>&1
              |grep "100% packet loss\|Host Unreachable"`
      if [ -z "$status" ] ;then
        # E3 is up again switch to best route
        route="best";
        ip route del 172.16.20.0/24 via 10.0.3.2 mpls $var_low
        ip route add 172.16.20.0/24 via 10.0.5.3 mpls $var_best
      fi
    ;;
  esac
done
```

We are on the best/medium route, ping to see if E3 is still up

E3 is down: traffic directed to network  
172.16.20.0/24 is switched from E3 to E2

# MPLS configuration on Ier E1

MPLS\_on\_E1.sh

```
route="best"
while [ "1" = "1" ]
do
  case $route in
    best)
      status=`/usr/sbin/hping3 -c 1 --icmp 10.0.5.3 2>&1
              |grep "100% packet loss\|Host Unreachable"`
      if [ ! -z "$status" ] ;then
        # E3 is down switch to low route
        route="low";
        ip route del 172.16.20.0/24 via 10.0.5.3 mpls $var_best
        ip route add 172.16.20.0/24 via 10.0.3.2 mpls $var_low
      fi
      ;;
    low)
      status=`/usr/sbin/hping3 -c 1 --icmp 10.0.5.3 2>&1
              |grep "100% packet loss\|Host Unreachable"`
      if [ -z "$status" ] ;then
        # E3 is up again switch to best route
        route="best";
        ip route del 172.16.20.0/24 via 10.0.3.2 mpls $var_low
        ip route add 172.16.20.0/24 via 10.0.5.3 mpls $var_best
      fi
      ;;
  esac
done
```

we are on the low route, ping to see if E3 is up

E3 is back up: traffic directed to network 172.16.20.0/24 is switched from E2 to E3



# MPLS configuration on Ier E1

- Note: the script for E1 collapses the *best* and *medium* routes into one case, because both routes reach E1 via the same link

# MPLS configuration on Ier E4

- Similar to E1
- The loop inside the scripts distinguishes among *best*, *medium*, and *low* routes by checking the status of link 6 and of E3 using ping packets

# MPLS configuration on Ixr E5

#A2->A3

#if input label is 100, keep it and push 10000 on top of it.

```
mpls labelspace set dev eth0 labelspace 0
```

```
mpls ilm add label gen 100 labelspace 0
```

```
key1=`mpls nhlf add key 0 instructions push gen 10000 nexthop  
eth1 ipv4 10.0.1.2 |grep key | cut -c 17-26`
```

```
key2=`mpls nhlf add key 0 instructions push gen 100 forward  
$key1 |grep key | cut -c 17-26`
```

With this instruction  
label 1000 is set on  
top of label 100

#This instruction exchange incoming label 100 with the label  
#stack (100, 10000)

```
mpls xc add ilm_label gen 100 ilm_labelspace 0 nhlf_key $key2
```

exchange

Incoming label

# MPLS configuration on Isr E5

```
#A3->A2
```

```
#if input label is 20000, pop it and keep the bottom label 600
```

```
mpls labelspace set dev eth1 labelspace 0
```

```
mpls ilm add label gen 20000 labelspace 0
```

```
mpls ilm add label gen 600 labelspace 0
```

```
key1=`mpls nhlfe add key 0 instructions push gen 600 nexthop  
eth0 ipv4 10.0.4.4 |grep key | cut -c 17-26`
```

```
mpls xc add ilm_label gen 600 ilm_labelspace 0 nhlfe_key $key1
```

Exchange...


...incoming label  
600 at the bottom  
of the stack (so it  
is removed from  
the stack)...

...with another label  
600

# MPLS configuration on other lsrs

- E2 and E3 are similar to E5. In addition to E5 they only have the loop instructions (very similar to the E1's):
  - E2 to check the status of E3
  - E3 to check the status of Link6
    - Side note: we simulate the fault of Link6 by bringing E3's eth1 down. Therefore, upon restoring this interface we need to also redefine the nhlfe

# Starting the lab



A terminal window titled "host machine" with standard window controls (minimize, maximize, close). The terminal shows the following commands and output:

```
user@localhost:~$ cd netkit-lab_MPLS_TE
user@localhost:~/netkit-lab_MPLS_TE$ lstart ||
```

- Currently selected routes are shown at the top of terminal windows of relevant routers

# MPLS traffic analysis

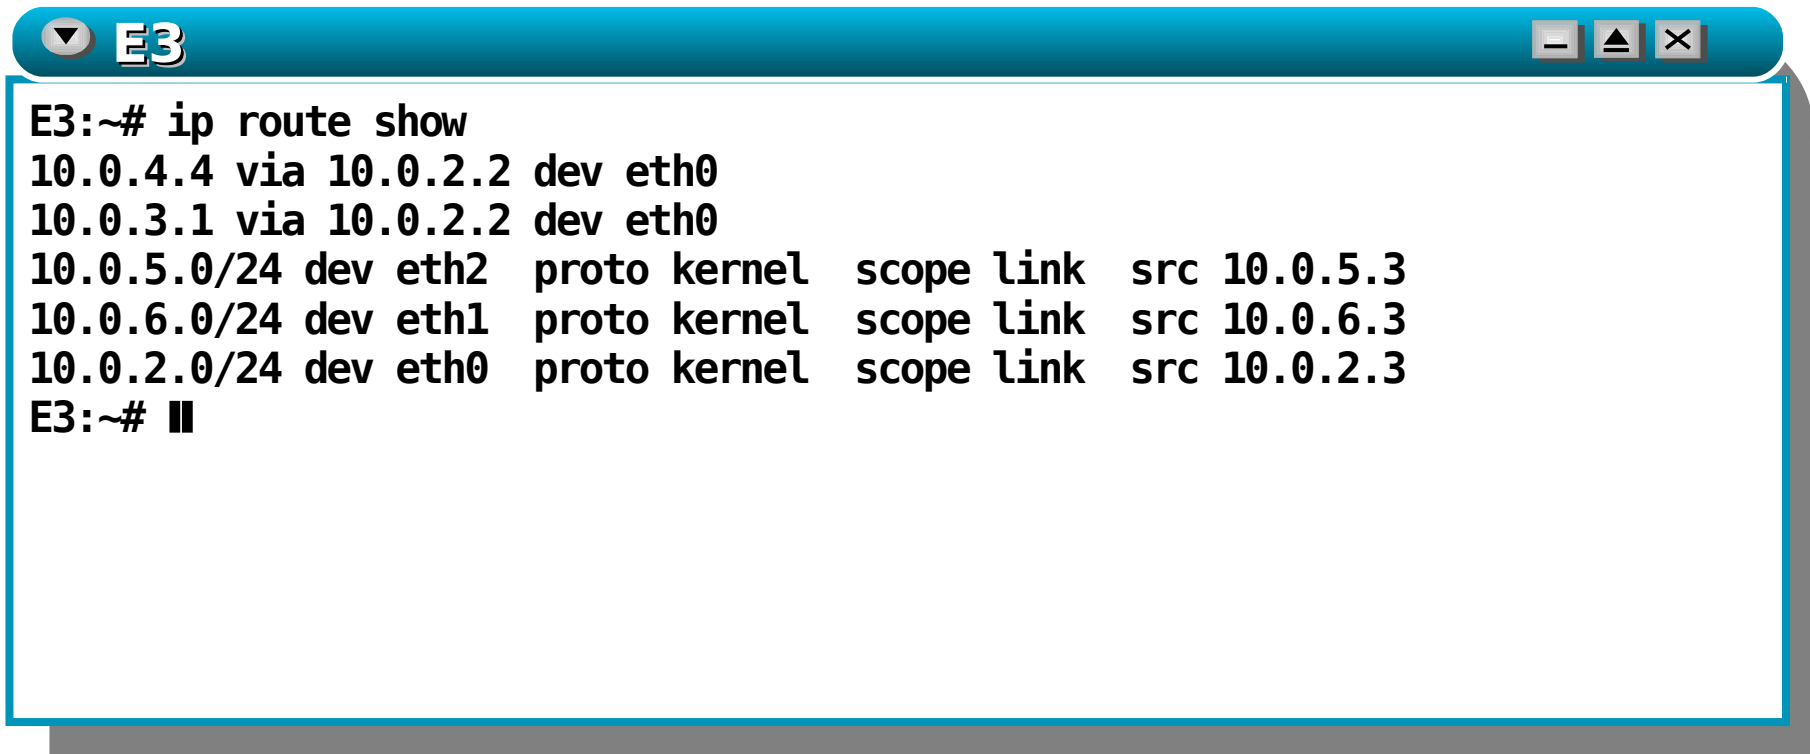
- Router E4 applies NHLFE 0x2 (= "push label 100") to traffic directed to 172.16.30.0/24 and forwards it to E3

▼ E4

```
E4:~# ip route show
10.0.2.3 via 10.0.4.5 dev eth0
10.0.4.0/24 dev eth0 proto kernel scope link src 10.0.4.4
172.16.20.0/24 dev eth2 proto kernel scope link src 172.16.20.4
10.0.6.0/24 dev eth1 proto kernel scope link src 10.0.6.4
172.16.30.0/24 via 10.0.6.3 dev eth1 mpls 0x2
E4:~# mpls nhlfe show
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
      push gen 100 set eth0 ipv4 10.0.4.5 (0 bytes, 0 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
      push gen 100 set eth1 ipv4 10.0.6.3 (8652 bytes, 103 pkts)
E4:~# █
```

# MPLS traffic analysis

- Isr E3's IP (correctly) knows nothing about 172.16.30.0/24 and 172.16.20.0/24

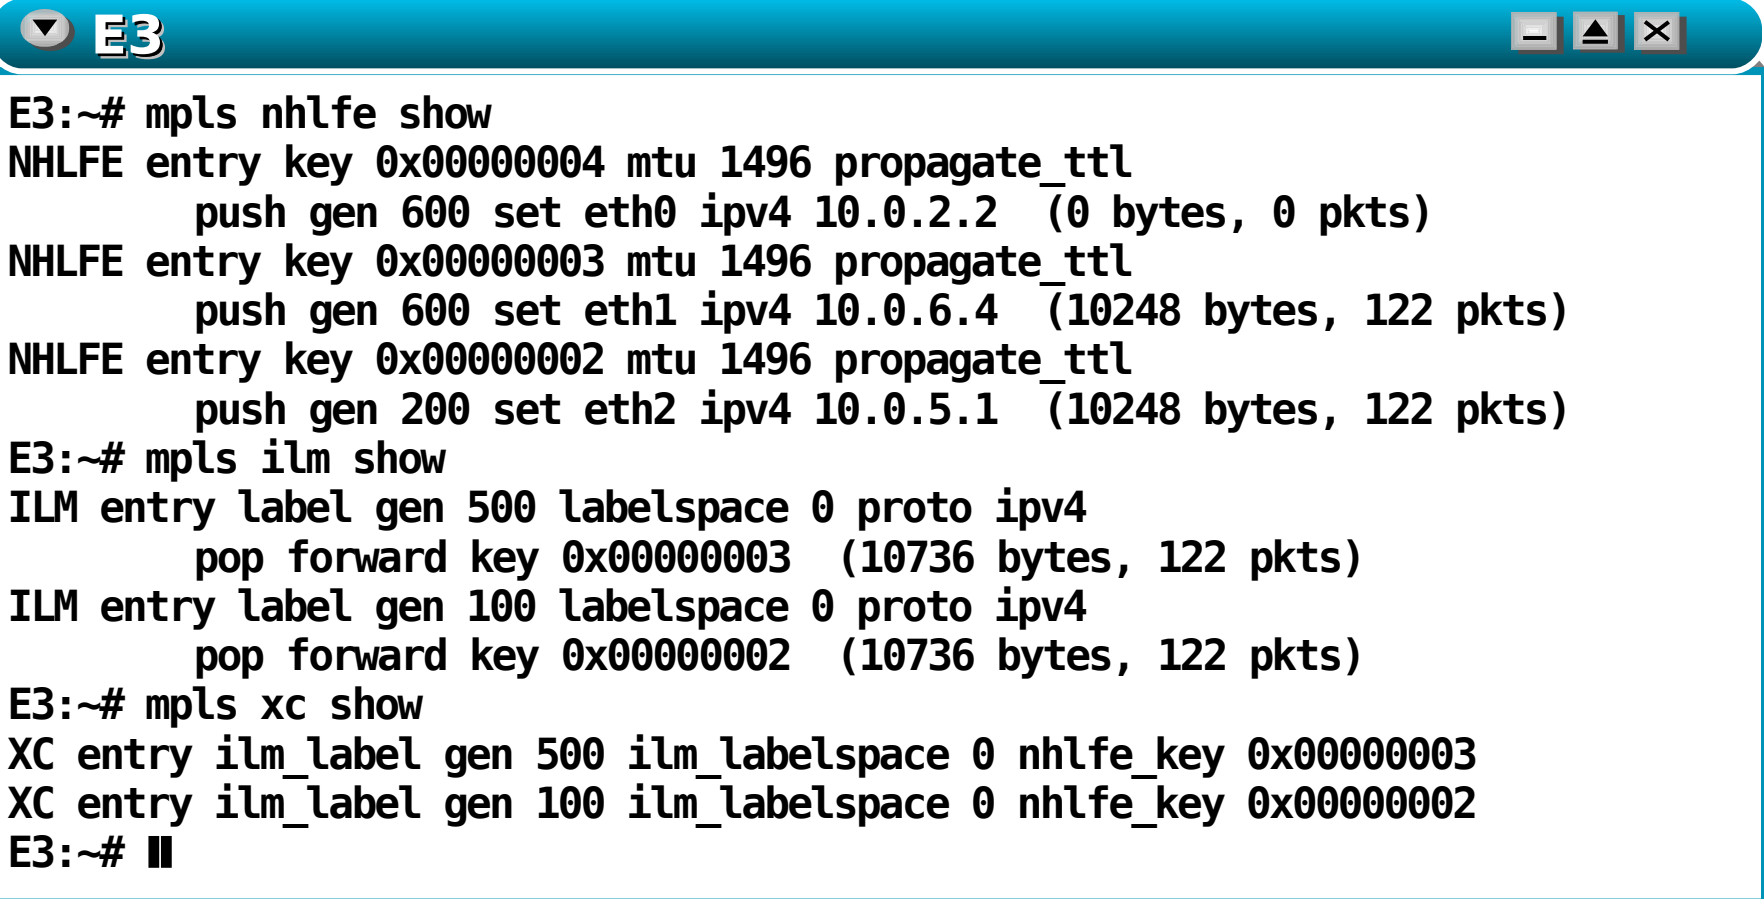


```
E3:~# ip route show
10.0.4.4 via 10.0.2.2 dev eth0
10.0.3.1 via 10.0.2.2 dev eth0
10.0.5.0/24 dev eth2 proto kernel scope link src 10.0.5.3
10.0.6.0/24 dev eth1 proto kernel scope link src 10.0.6.3
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.3
E3:~# █
```



# MPLS traffic analysis

- Isr E3's MPLS knows how to forward labeled packets



```
E3:~# mpls nhlfe show
NHLFE entry key 0x00000004 mtu 1496 propagate_ttl
    push gen 600 set eth0 ipv4 10.0.2.2 (0 bytes, 0 pkts)
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 600 set eth1 ipv4 10.0.6.4 (10248 bytes, 122 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
    push gen 200 set eth2 ipv4 10.0.5.1 (10248 bytes, 122 pkts)
E3:~# mpls ilm show
ILM entry label gen 500 labelspace 0 proto ipv4
    pop forward key 0x00000003 (10736 bytes, 122 pkts)
ILM entry label gen 100 labelspace 0 proto ipv4
    pop forward key 0x00000002 (10736 bytes, 122 pkts)
E3:~# mpls xc show
XC entry ilm_label gen 500 ilm_labelspace 0 nhlfe_key 0x00000003
XC entry ilm_label gen 100 ilm_labelspace 0 nhlfe_key 0x00000002
E3:~# █
```

# MPLS traffic analysis

- incoming packets with label 100 have their label recognized and popped...

```
E3:~# mpls nhlfe show
NHLFE entry key 0x00000004 mtu 1496 propagate_ttl
    push gen 600 set eth0 ipv4 10.0.2.2 (0 bytes, 0 pkts)
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 600 set eth1 ipv4 10.0.6.4 (10248 bytes, 122 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
    push gen 200 set eth2 ipv4 10.0.5.1 (10248 bytes, 122 pkts)
E3:~# mpls ilm show
ILM entry label gen 500 labelspace 0 proto ipv4
    pop forward key 0x00000003 (10736 bytes, 122 pkts)
ILM entry label gen 100 labelspace 0 proto ipv4
    pop forward key 0x00000002 (10736 bytes, 122 pkts)
E3:~# mpls xc show
XC entry ilm_label gen 500 ilm_labelspace 0 nhlfe_key 0x00000003
XC entry ilm_label gen 100 ilm_labelspace 0 nhlfe_key 0x00000002
E3:~# █
```

# MPLS traffic analysis


- ...and are forwarded to E1 after swapping the label with 200

```
E3:~# mpls nhlfe show
NHLFE entry key 0x00000004 mtu 1496 propagate_ttl
    push gen 600 set eth0 ipv4 10.0.2.2 (0 bytes, 0 pkts)
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 600 set eth1 ipv4 10.0.6.4 (10248 bytes, 122 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
    push gen 200 set eth2 ipv4 10.0.5.1 (10248 bytes, 122 pkts)
E3:~# mpls ilm show
ILM entry label gen 500 labelspace 0 proto ipv4
    pop forward key 0x00000003 (10736 bytes, 122 pkts)
ILM entry label gen 100 labelspace 0 proto ipv4
    pop forward key 0x00000002 (10736 bytes, 122 pkts)
E3:~# mpls xc show
XC entry ilm_label gen 500 ilm_labelspace 0 nhlfe_key 0x00000003
XC entry ilm_label gen 100 ilm_labelspace 0 nhlfe_key 0x00000002
E3:~# █
```

# MPLS traffic analysis

- We see how MPLS fills its tables and the integration between IP and MPLS
- You can try, for exercise, to see other router's and LSR's routing tables to check their correctness

# Normal traffic



# MPLS Traffic analysis

- Some checkpoints have been identified along the path from A2 to A3
  - E4: both interfaces, to observe traffic before and after the insertion of the MPLS header
  - E3: interface eth2 to observe label switching
  - E1: interface eth1, where the MPLS header is removed
- Similar checkpoints can be considered for the traffic from A3 to A2

# MPLS Traffic analisys

- We sniff packets using tcpdump and examine the dumps on the host using wireshark

▼ A2

```
A2:~# ping 172.16.30.30 ||
```

▼ EX

```
EX:~# tcpdump -i ethY -w /hostlab/sniff_${HOSTNAME}_Y.cap -s 1500 ||
```

▼ host machine

```
user@localhost:~/netkit-lab_MPLS_TE$ wireshark -r sniff_EX_Y.cap ||
```

# MPLS Traffic analysis

- Sniffing on E4's eth2: plain ICMP packets

No.	Time	Source	Destination	Protocol	Info
1	14:40:01.452264	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request (id
2	14:40:01.452785	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply (id
3	14:40:02.451335	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request (id
4	14:40:02.451952	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply (id
5	14:40:03.450264	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request (id
6	14:40:03.450655	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply (id
7	14:40:04.449304	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request (id
8	14:40:04.449728	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply (id



# MPLS Traffic analysis

- Sniffing on E4's eth1: ICMP packets encapsulated in MPLS

No..	Time	Source	Destination	Protocol	Info
1	0.000000	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
2	0.000037	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
3	0.078153	10.0.6.4	10.0.6.3	ICMP	Echo (ping) request
4	0.079557	10.0.6.3	10.0.6.4	ICMP	Echo (ping) reply
5	0.238018	10.0.6.4	10.0.6.3	ICMP	Echo (ping) request
6	0.238353	10.0.6.3	10.0.6.4	ICMP	Echo (ping) reply
7	0.398667	10.0.6.4	10.0.6.3	ICMP	Echo (ping) request
8	0.399163	10.0.6.3	10.0.6.4	ICMP	Echo (ping) reply

.....

▶ Frame 1 (102 bytes on wire, 102 bytes captured)

▶ Ethernet II, Src: 62:bc:16:f4:cc:7c (62:bc:16:f4:cc:7c), Dst: 42:dd:4f:2b:72:80 (42:dd:4f:2b:72:80)

▼ MultiProtocol Label Switching Header, Label: 100, Exp: 0, S: 1, TTL: 63

MPLS Label: 100

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 63

▶ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▶ Internet Control Message Protocol

Pushed label 100

# MPLS Traffic analysis

- Sniffing on E4's eth1: ICMP packets encapsulated in MPLS

No..	Time	Source	Destination	Protocol	Info
1	0.000000	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
2	0.000037	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
3	0.078153	10.0.6.4	10.0.6.3	ICMP	Echo (ping) request
4	0.079557	10.0.6.3	10.0.6.4	ICMP	Echo (ping) reply
5	0.238018	10.0.6.4	10.0.6.3	ICMP	Echo (ping) request
6	0.238353	10.0.6.3	10.0.6.4	ICMP	Echo (ping) reply
7	0.398667	10.0.6.4	10.0.6.3	ICMP	Echo (ping) request
8	0.399163	10.0.6.3	10.0.6.4	ICMP	Echo (ping) reply

▷ Frame 2 (102 bytes on wire, 102 bytes captured)

▷ Ethernet II, Src: 42:dd:4f:2b:72:80 (42:dd:4f:2b:72:80), Dst: 62:bc:16:f4:cc:7c (62:bc:16:f4:cc:7c)

▽ MultiProtocol Label Switching Header, Label: 600, Exp: 0, S: 1, TTL: 62

MPLS Label: 600

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 62

Received label 600

▷ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)

▷ Internet Control Message Protocol

# MPLS Traffic analysis

- Sniffing on E3's eth2: ICMP packets encapsulated in MPLS with a label swapped

No. .	Time	Source	Destination	Protocol	Info
11	0.754606	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
12	0.758501	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
13	0.786299	10.0.5.1	10.0.5.3	ICMP	Echo (ping) request
14	0.786324	10.0.5.3	10.0.5.1	ICMP	Echo (ping) reply
15	0.966331	10.0.5.1	10.0.5.3	ICMP	Echo (ping) request
16	0.966365	10.0.5.3	10.0.5.1	ICMP	Echo (ping) reply

.....

▷ Frame 11 (102 bytes on wire, 102 bytes captured)

▷ Ethernet II, Src: a2:12:57:55:dd:d0 (a2:12:57:55:dd:d0), Dst: 9e:8c:07:a2:c3:d6 (9e:8c:07:a2:c3:d6)

▽ MultiProtocol Label Switching Header, Label: 200, Exp: 0, S: 1, TTL: 62

MPLS Label: 200

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 62

▷ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▷ Internet Control Message Protocol

Pushed label 200

# MPLS Traffic analysis

## ■ Sniffing on E3's eth2

No. .	Time	Source	Destination	Protocol	Info
11	0.754606	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
12	0.758501	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
13	0.786299	10.0.5.1	10.0.5.3	ICMP	Echo (ping) request
14	0.786324	10.0.5.3	10.0.5.1	ICMP	Echo (ping) reply
15	0.966331	10.0.5.1	10.0.5.3	ICMP	Echo (ping) request
16	0.966365	10.0.5.3	10.0.5.1	ICMP	Echo (ping) reply

.....

▶ Frame 12 (102 bytes on wire, 102 bytes captured)

▶ Ethernet II, Src: 9e:8c:07:a2:c3:d6 (9e:8c:07:a2:c3:d6), Dst: a2:12:57:55:dd:d0 (a2:12:57:55:dd:d0)

▼ MultiProtocol Label Switching Header, Label: 500, Exp: 0, S: 1, TTL: 63

MPLS Label: 500

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 63

▶ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)

▶ Internet Control Message Protocol

Received label 500

# MPLS Traffic analysis

- Sniffing on E1's eth1: (back to) plain ICMP packets

No..	Time	Source	Destination	Protocol	Info
1	0.000000	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
2	0.000084	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
3	1.007343	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
4	1.007475	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
5	2.010463	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
6	2.010657	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply

.....

▶ Frame 1 (98 bytes on wire, 98 bytes captured)

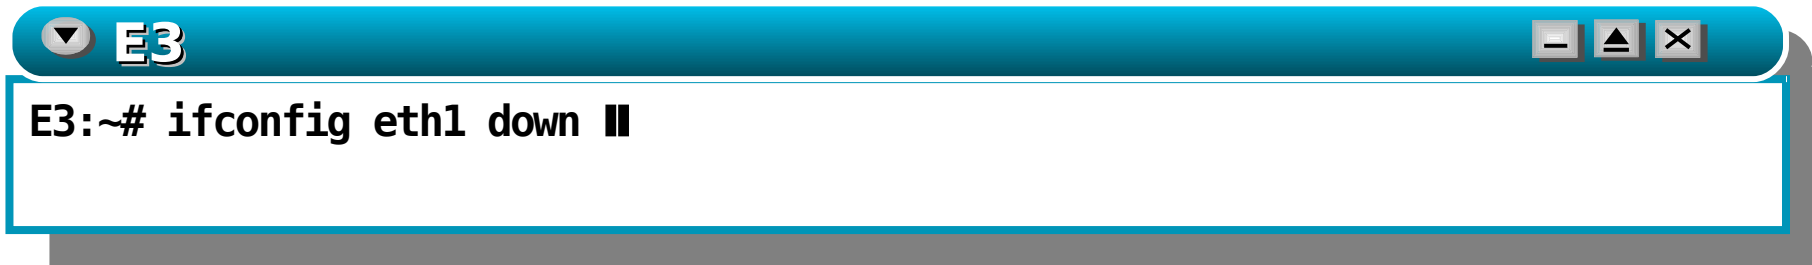
▶ Ethernet II, Src: be:af:f2:d2:70:1e (be:af:f2:d2:70:1e), Dst: 5e:a2:21:c6:8d:46 (5e:a2:21:c6:8d:46)

▶ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▶ Internet Control Message Protocol


# MPLS Traffic analysis

- Let's try to put link 6 down



- Routes taken by E1-E4 are updated (look at the currently selected routes in their windows)
- Note: due to the time required for detection, a few ICMP packets may be lost (or enqueued at some interface until it is brought back up)

# Traffic with link 6 down



# MPLS traffic analysis

## ■ Sniffing on E4's eth0

No..	Time	Source	Destination	Protocol	Info
5	0.812049	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
6	0.812864	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
7	0.961206	10.0.4.4	10.0.2.3	ICMP	Echo (ping) request
8	1.822025	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
9	1.828680	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
10	2.832026	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request

- ▶ Frame 5 (102 bytes on wire, 102 bytes captured)
- ▶ Ethernet II, Src: 96:9d:56:2b:d6:68 (96:9d:56:2b:d6:68), Dst: 0a:b7:2c:3a:60:63 (0a:b7:2c:3a:60:63)
- ▼ MultiProtocol Label Switching Header, Label: 100, Exp: 0, S: 1, TTL: 63
- MPLS Label: 100
  - MPLS Experimental Bits: 0
  - MPLS Bottom Of Label Stack: 1
  - MPLS TTL: 63
- ▶ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)
- ▶ Internet Control Message Protocol

Pushed label 100



# MPLS traffic analysis

## ■ Sniffing on E4's eth0

No..	Time	Source	Destination	Protocol	Info
5	0.812049	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
6	0.812864	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
7	0.961206	10.0.4.4	10.0.2.3	ICMP	Echo (ping) request
8	1.822025	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
9	1.828680	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
10	2.832026	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request

- ▶ Frame 6 (102 bytes on wire, 102 bytes captured)
- ▶ Ethernet II, Src: 0a:b7:2c:3a:60:63 (0a:b7:2c:3a:60:63), Dst: 96:9d:56:2b:d6:68 (96:9d:56:2b:d6:68)
- ▼ MultiProtocol Label Switching Header, Label: 600, Exp: 0, S: 1, TTL: 60
  - MPLS Label: 600
  - MPLS Experimental Bits: 0
  - MPLS Bottom Of Label Stack: 1
  - MPLS TTL: 60
- ▶ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)
- ▶ Internet Control Message Protocol

Received label 600

# MPLS traffic analysis

## ■ Sniffing on E5's eth1

No..	Time	Source	Destination	Protocol	Info
6	0.976778	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
7	0.983217	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
8	1.976253	10.0.4.4	10.0.2.3	ICMP	Echo (ping) request
9	1.987727	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
10	1.996001	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply

▶ Frame 6 (106 bytes on wire, 106 bytes captured)

▶ Ethernet II, Src: 4a:7c:7f:88:47:e5 (4a:7c:7f:88:47:e5), Dst: ba:25:23:ed:66:84 (ba:25:23:ed:66:84)

▼ MultiProtocol Label Switching Header, Label: 10000, Exp: 0, S: 0, TTL: 62

MPLS Label: 10000

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 0

MPLS TTL: 62

External label 10000

▼ MultiProtocol Label Switching Header, Label: 100, Exp: 0, S: 1, TTL: 62

MPLS Label: 100

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 62

Internal label 100

▶ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▶ Internet Control Message Protocol

# MPLS traffic analysis

## ■ Sniffing on E5's eth1

No..	Time	Source	Destination	Protocol	Info
6	0.976778	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
7	0.983217	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
8	1.976253	10.0.4.4	10.0.2.3	ICMP	Echo (ping) request
9	1.987727	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
10	1.996001	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply

▷ Frame 7 (106 bytes on wire, 106 bytes captured)

▷ Ethernet II, Src: ba:25:23:ed:66:84 (ba:25:23:ed:66:84), Dst: 4a:7c:7f:88:47:e5 (4a:7c:7f:88:47:e5)

▽ MultiProtocol Label Switching Header, Label: 20000, Exp: 0, S: 0, TTL: 61

MPLS Label: 20000

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 0

MPLS TTL: 61

Receive internal label 20000

▽ MultiProtocol Label Switching Header, Label: 600, Exp: 0, S: 1, TTL: 61

MPLS Label: 600

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 61

Received external label 600

▷ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)

▷ Internet Control Message Protocol

# MPLS traffic analysis

## ■ Sniffing on E2's eth0

No..	Time	Source	Destination	Protocol	Info
24	2.110859	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
25	2.111568	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
26	2.169671	10.0.2.2	10.0.2.3	ICMP	Echo (ping) request
27	2.170013	10.0.2.3	10.0.2.2	ICMP	Echo (ping) reply
28	2.329784	10.0.2.2	10.0.2.3	ICMP	Echo (ping) request

.....

▶ Frame 24 (102 bytes on wire, 102 bytes captured)

▶ Ethernet II, Src: 0a:6d:10:7d:9f:fc (0a:6d:10:7d:9f:fc), Dst: 16:03:42:67:ad:a3 (16:03:42:67:ad:a3)

▼ MultiProtocol Label Switching Header, Label: 100, Exp: 0, S: 1, TTL: 61

MPLS Label: 100

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 61

▶ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▶ Internet Control Message Protocol

Pushed label 100

# MPLS traffic analysis

## ■ Sniffing on E2's eth0

No..	Time	Source	Destination	Protocol	Info
24	2.110859	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
25	2.111568	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
26	2.169671	10.0.2.2	10.0.2.3	ICMP	Echo (ping) request
27	2.170013	10.0.2.3	10.0.2.2	ICMP	Echo (ping) reply
28	2.329784	10.0.2.2	10.0.2.3	ICMP	Echo (ping) request

.....

▶ Frame 25 (102 bytes on wire, 102 bytes captured)

▶ Ethernet II, Src: 16:03:42:67:ad:a3 (16:03:42:67:ad:a3), Dst: 0a:6d:10:7d:9f:fc (0a:6d:10:7d:9f:fc)

▼ MultiProtocol Label Switching Header, Label: 600, Exp: 0, S: 1, TTL: 62

    MPLS Label: 600

    MPLS Experimental Bits: 0

    MPLS Bottom Of Label Stack: 1

    MPLS TTL: 62

▶ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)

▶ Internet Control Message Protocol

Received label 600

# MPLS traffic analysis

## ■ Sniffing on E3's eth2

No.	Time	Source	Destination	Protocol	Info
5	0.253065	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
6	0.257714	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
7	0.305738	10.0.5.1	10.0.5.3	ICMP	Echo (ping) request
8	0.305772	10.0.5.3	10.0.5.1	ICMP	Echo (ping) reply
Frame 5 (102 bytes on wire, 102 bytes captured)					
Ethernet II, Src: a2:12:57:55:dd:d0 (a2:12:57:55:dd:d0), Dst: 9e:8c:07:a2:c3:d6 (9e:8c:07:a2:c3:d6)					
MultiProtocol Label Switching Header, Label: 200, Exp: 0, S: 1, TTL: 60					
MPLS Label: 200					
MPLS Experimental Bits: 0					
MPLS Bottom Of Label Stack: 1					
MPLS TTL: 60					
Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)					
Internet Control Message Protocol					

Pushed label 200

# MPLS traffic analysis

## ■ Sniffing on E3's eth2

No.	Time	Source	Destination	Protocol	Info
5	0.253065	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
6	0.257714	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
7	0.305738	10.0.5.1	10.0.5.3	ICMP	Echo (ping) request
8	0.305772	10.0.5.3	10.0.5.1	ICMP	Echo (ping) reply
9	0.455001	10.0.5.1	10.0.5.3	ICMP	Echo (ping) request

▶ Frame 6 (102 bytes on wire, 102 bytes captured)

▶ Ethernet II, Src: 9e:8c:07:a2:c3:d6 (9e:8c:07:a2:c3:d6), Dst: a2:12:57:55:dd:d0 (a2:12:57:55:dd:d0)

▼ MultiProtocol Label Switching Header, Label: 500, Exp: 0, S: 1, TTL: 63

- MPLS Label: 500
- MPLS Experimental Bits: 0
- MPLS Bottom Of Label Stack: 1
- MPLS TTL: 63

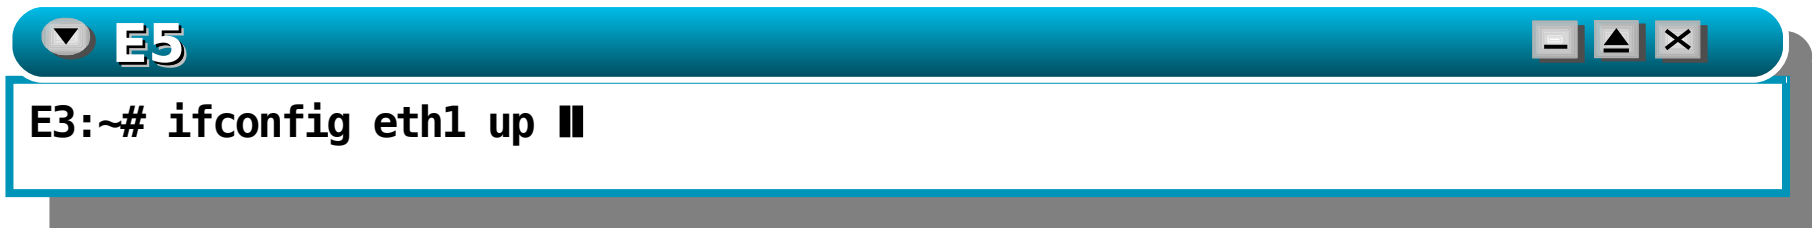
▶ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)

▶ Internet Control Message Protocol

Received label 500

# MPLS Traffic analysis

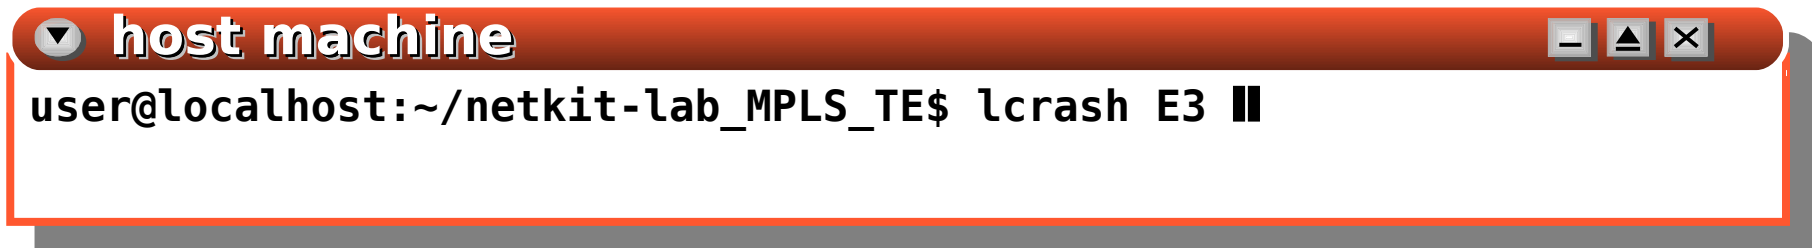
- We bring link 6 up again, you can check that traffic will switch to best



A terminal window with a blue title bar labeled 'E5'. The terminal text shows the command 'ifconfig eth1 up' being entered at the prompt 'E3:~#'. The window has standard minimize, maximize, and close buttons on the right.

```
E3:~# ifconfig eth1 up ||
```

- Now we bring E3 down and see what happens




A terminal window with an orange title bar labeled 'host machine'. The terminal text shows the command 'lcrash E3' being entered at the prompt 'user@localhost:~/netkit-lab\_MPLS\_TE\$'. The window has standard minimize, maximize, and close buttons on the right.

```
user@localhost:~/netkit-lab_MPLS_TE$ lcrash E3 ||
```



# Traffic with E3 down



# MPLS traffic analysis

## ■ Sniffing on E4's eth0

No..	Time	Source	Destination	Protocol	Info
3	0.743103	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
4	0.743587	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
5	0.755362	10.0.4.4	10.0.2.3	ICMP	Echo (ping) request
6	1.743167	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
7	1.743860	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply

▶ Frame 3 (102 bytes on wire, 102 bytes captured)

▶ Ethernet II, Src: 96:9d:56:2b:d6:68 (96:9d:56:2b:d6:68), Dst: 0a:b7:2c:3a:60:63 (0a:b7:2c:3a:60:63)

▼ MultiProtocol Label Switching Header, Label: 100, Exp: 0, S: 1, TTL: 63

- MPLS Label: 100
- MPLS Experimental Bits: 0
- MPLS Bottom Of Label Stack: 1
- MPLS TTL: 63

▶ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▶ Internet Control Message Protocol

Pushed label 100

# MPLS traffic analysis

## ■ Sniffing on E4's eth0

No.	Time	Source	Destination	Protocol	Info
3	0.743103	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
4	0.743587	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
5	0.755362	10.0.4.4	10.0.2.3	ICMP	Echo (ping) request
6	1.743167	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
7	1.743860	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply

- ▶ Frame 4 (102 bytes on wire, 102 bytes captured)
- ▶ Ethernet II, Src: 0a:b7:2c:3a:60:63 (0a:b7:2c:3a:60:63), Dst: 96:9d:56:2b:d6:68 (96:9d:56:2b:d6:68)
- ▼ MultiProtocol Label Switching Header, Label: 600, Exp: 0, S: 1, TTL: 61
- MPLS Label: 600
  - MPLS Experimental Bits: 0
  - MPLS Bottom Of Label Stack: 1
  - MPLS TTL: 61
- ▶ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)
- ▶ Internet Control Message Protocol

Received label 600

# MPLS traffic analysis

## ■ Sniffing on E5's eth1

No..	Time	Source	Destination	Protocol	Info
2	0.239797	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
3	0.240273	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
4	1.249740	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
5	1.250294	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
6	2.259664	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request

▷ Frame 2 (106 bytes on wire, 106 bytes captured)

▷ Ethernet II, Src: 4a:7c:7f:88:47:e5 (4a:7c:7f:88:47:e5), Dst: ba:25:23:ed:66:84 (ba:25:23:ed:66:84)

▽ MultiProtocol Label Switching Header, Label: 10000, Exp: 0, S: 0, TTL: 62

MPLS Label: 10000

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 0

MPLS TTL: 62

External label 10000

▽ MultiProtocol Label Switching Header, Label: 100, Exp: 0, S: 1, TTL: 62

MPLS Label: 100

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 62

Internal label 100

▷ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▷ Internet Control Message Protocol

# MPLS traffic analysis

## ■ Sniffing on E5's eth1

No..	Time	Source	Destination	Protocol	Info
2	0.239797	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
3	0.240273	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
4	1.249740	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
5	1.250294	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
6	2.259664	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request

▷ Frame 3 (106 bytes on wire, 106 bytes captured)

▷ Ethernet II, Src: ba:25:23:ed:66:84 (ba:25:23:ed:66:84), Dst: 4a:7c:7f:88:47:e5 (4a:7c:7f:88:47:e5)

▼ MultiProtocol Label Switching Header, Label: 20000, Exp: 0, S: 0, TTL: 62

MPLS Label: 20000

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 0

MPLS TTL: 62

Received external label 20000

▼ MultiProtocol Label Switching Header, Label: 600, Exp: 0, S: 1, TTL: 62

MPLS Label: 600

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 62

Received internal label 600

▷ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)

▷ Internet Control Message Protocol

# MPLS traffic analysis

## ■ Sniffing on E2's eth2

No..	Time	Source	Destination	Protocol	Info
1	0.000000	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
2	0.000293	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
3	1.009738	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
4	1.010021	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
5	2.009940	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request

.....

▶ Frame 1 (102 bytes on wire, 102 bytes captured)

▶ Ethernet II, Src: ce:97:71:1f:f6:d1 (ce:97:71:1f:f6:d1), Dst: 66:8b:cf:da:45:86 (66:8b:cf:da:45:86)

▼ MultiProtocol Label Switching Header, Label: 200, Exp: 0, S: 1, TTL: 61

MPLS Label: 200

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 61

▶ Internet Protocol, Src: 172.16.20.20 (172.16.20.20), Dst: 172.16.30.30 (172.16.30.30)

▶ Internet Control Message Protocol

Pushed label 200

# MPLS traffic analysis

## ■ Sniffing on E2's eth2

No..	Time	Source	Destination	Protocol	Info
1	0.000000	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
2	0.000293	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
3	1.009738	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request
4	1.010021	172.16.30.30	172.16.20.20	ICMP	Echo (ping) reply
5	2.009940	172.16.20.20	172.16.30.30	ICMP	Echo (ping) request

▷ Frame 2 (102 bytes on wire, 102 bytes captured)

▷ Ethernet II, Src: 66:8b:cf:da:45:86 (66:8b:cf:da:45:86), Dst: ce:97:71:1f:f6:d1 (ce:97:71:1f:f6:d1)

▼ MultiProtocol Label Switching Header, Label: 600, Exp: 0, S: 1, TTL: 63

MPLS Label: 600

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

MPLS TTL: 63

▷ Internet Protocol, Src: 172.16.30.30 (172.16.30.30), Dst: 172.16.20.20 (172.16.20.20)

▷ Internet Control Message Protocol



Received label 600

# Proposed exercises

- Sniff on all the interfaces with wireshark
- Try to change NHLFE tables on the running lab, adding or deleting entries (use “mpls -help” inside virtual machines to see the help)
- Try to change label numbers and encapsulation