

netkit lab

MPLS for Linux

Version	1.1
Author(s)	S.Filippi, F.Antonini (Kasko Networks S.r.l.) Reviewed by M. Rimondini (Roma Tre Univ.)
E-mail	silvia.filippi@kaskonetworks.it
Web	http://www.kaskonetworks.it/
Description	An example with 2 lers, 1 lsr and 3 hosts to show how mpls works on ipv4

copyright notice

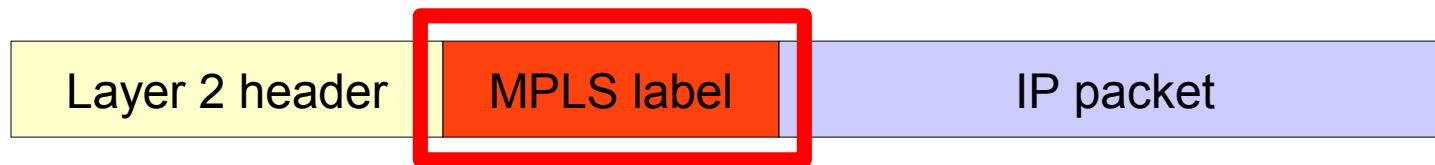
- All the pages/slides in this presentation, including but not limited to, images, photos, animations, videos, sounds, music, and text (hereby referred to as “material”) are protected by copyright.
- This material, with the exception of some multimedia elements licensed by other organizations, is property of the authors and/or organizations appearing in the first slide.
- This material, or its parts, can be reproduced and used for didactical purposes within universities and schools, provided that this happens for non-profit purposes.
- Information contained in this material cannot be used within network design projects or other products of any kind.
- Any other use is prohibited, unless explicitly authorized by the authors on the basis of an explicit agreement.
- The authors assume no responsibility about this material and provide this material “as is”, with no implicit or explicit warranty about the correctness and completeness of its contents, which may be subject to changes.
- This copyright notice must always be redistributed together with the material, or its portions.

MPLS basics (1/3)

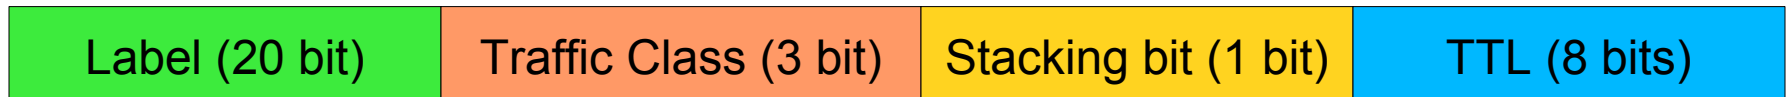
- MPLS stands for Multi Protocol Label Switching
- MPLS forwards packets inside the network using the “label switching” mechanism
- The mechanism is as follows:
 - Traffic that is supposed to be forwarded inside an MPLS network is classified (e.g., all the packets having the same destination address)
 - Each classified packet is assigned a label ([label binding](#)) and injected inside the MPLS network
 - Forwarding inside the MPLS network happens by just [swapping](#) the label (i.e., no more “routing by network address”)
 - “swapping”=replacing the label with another (possibly different) one; the scope of each label is a single link
 - The label is removed at the egress from the MPLS network

MPLS basics (2/3)

- After label insertion, the routed packet will look like this...



...with the *MPLS label* consisting of these fields:



- **Label**: the label value
- **Traffic Class**: used to differentiate priority among different types of traffic
- **Stacking bit**: a packet may be assigned a stack of labels, instead of a single one; this bit marks the end of the stack
- **TTL**: the usual Time to Live

MPLS basics (3/3)

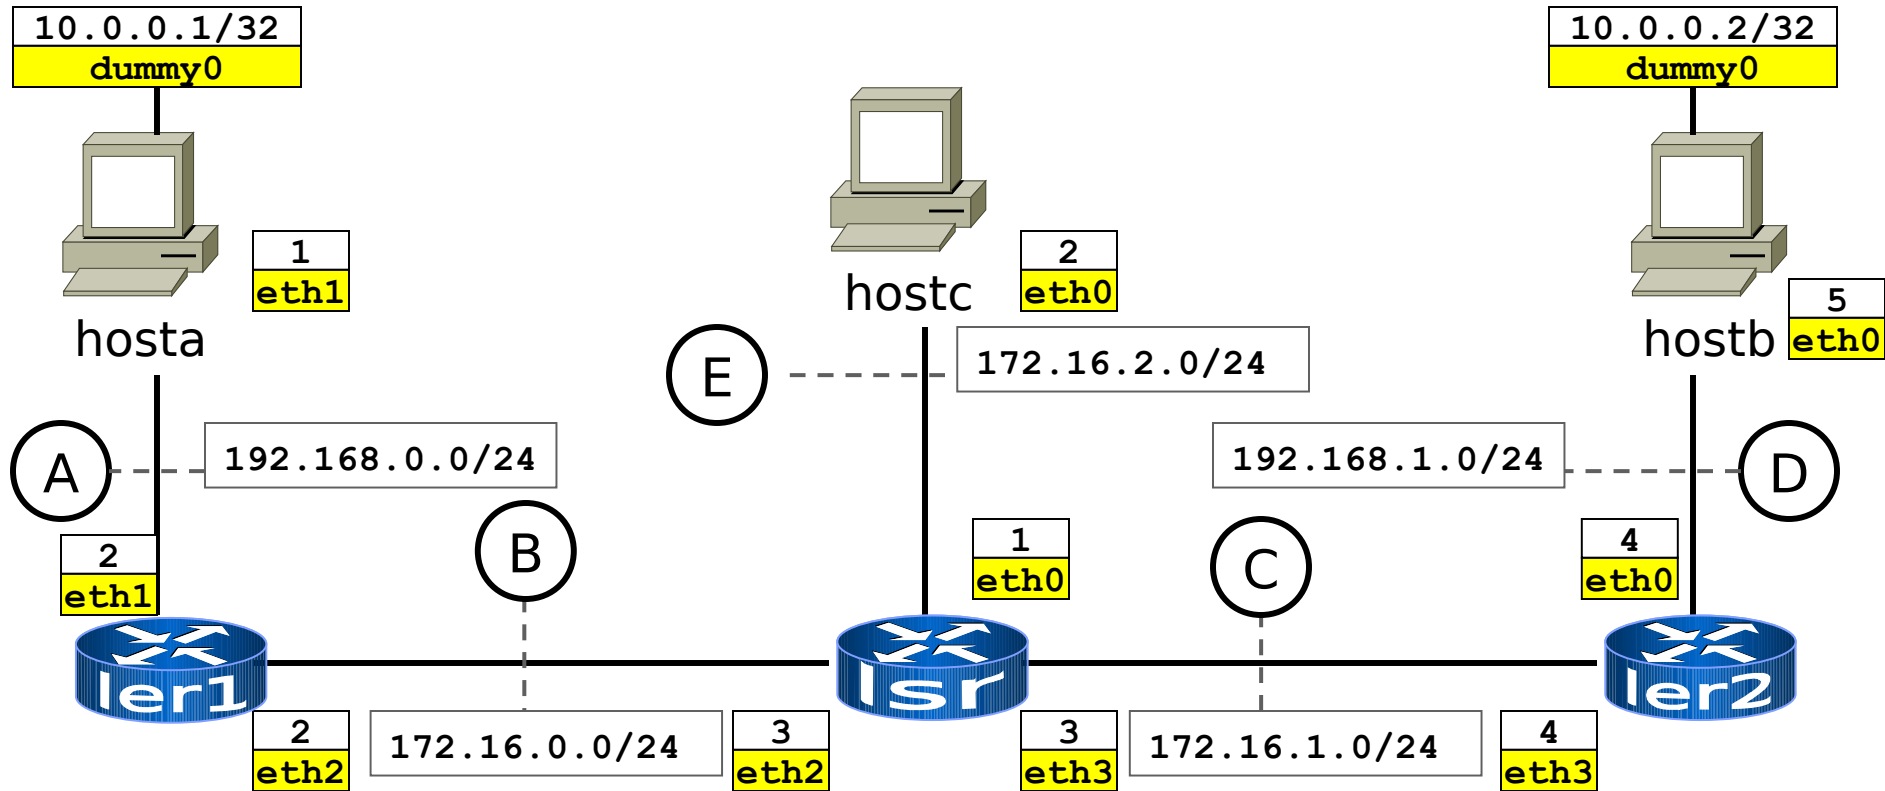
- In order to perform packet forwarding, MPLS routers will check some tables:
 - **NHLFE** (Next Hop Label Forwarding Entry): contains information about how to forward a packet in MPLS, including:
 - actual IP next hop for the packet
 - operations (push/pop) to be performed on the label stack
 - **ILM** (Incoming Label Map): maps incoming labelled packets to a NHLFE, i.e., tells how to forward labelled packets
 - **XC** (Cross Connect): maps an entry in the ILM to a NHLFE; tells the router how to perform label swapping

The MPLS Lab

- This lab is (highly) inspired to the example “*MPLS for Linux: IPv4 over MPLS: two LER one LSR example for mpls-linux-1.95x*” from [1]
- Some more MPLS terminology:
 - **LER** (Label Edge Router): entry and exit routers for MPLS networks. They push labels in packets entering the MPLS network and pop labels from packets that exit the MPLS network.
 - **LSR** (Label Switching Router): performs routing inside an MPLS network, based on solely swapping labels.
 - **FEC** (Forwarding Equivalence Class): a set of IP packets which are forwarded in the same manner

[1] <http://sourceforge.net/apps/mediawiki/mpls-linux/index.php?title=Examples>

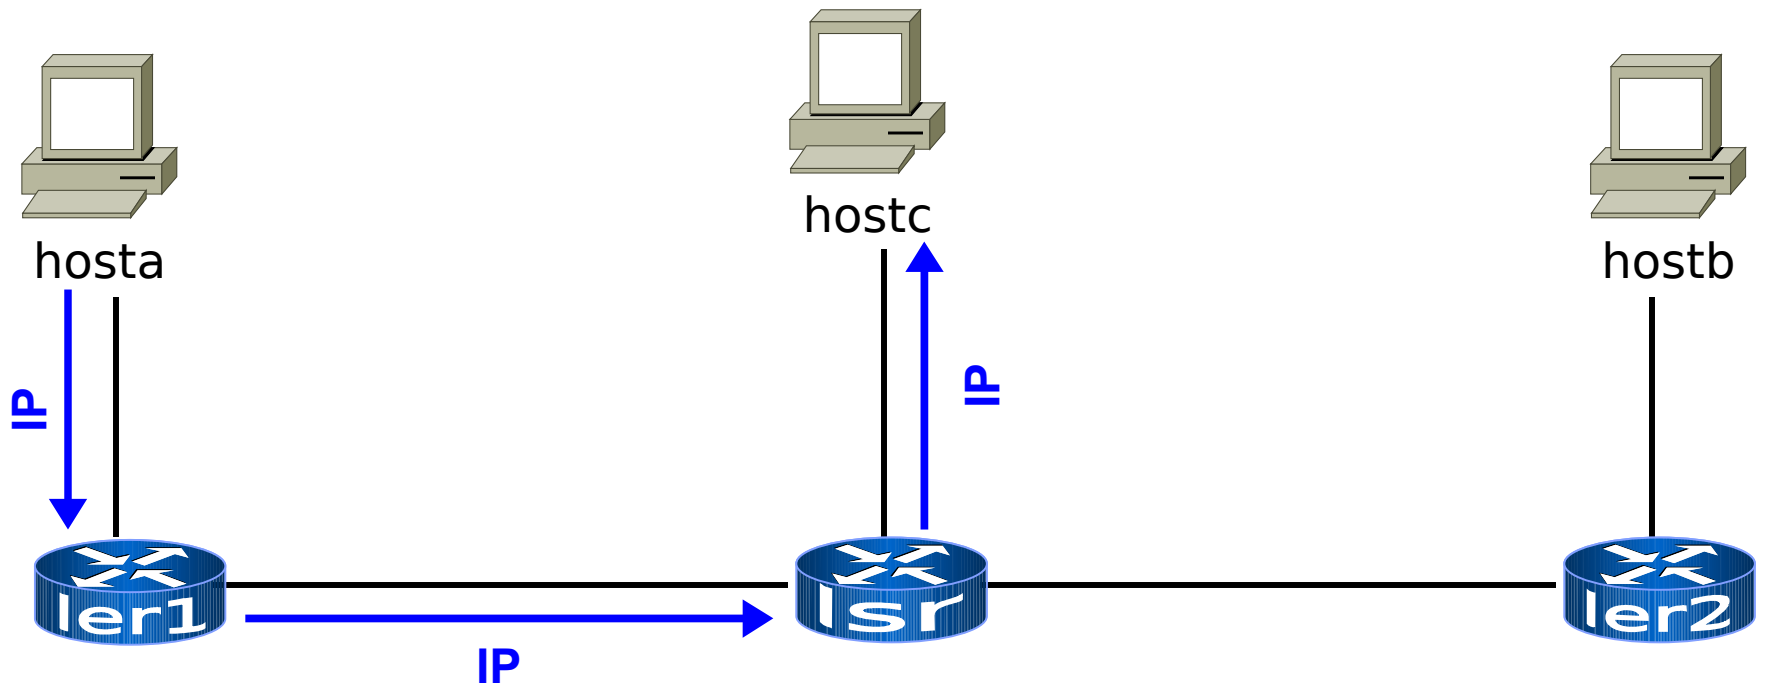
Topology



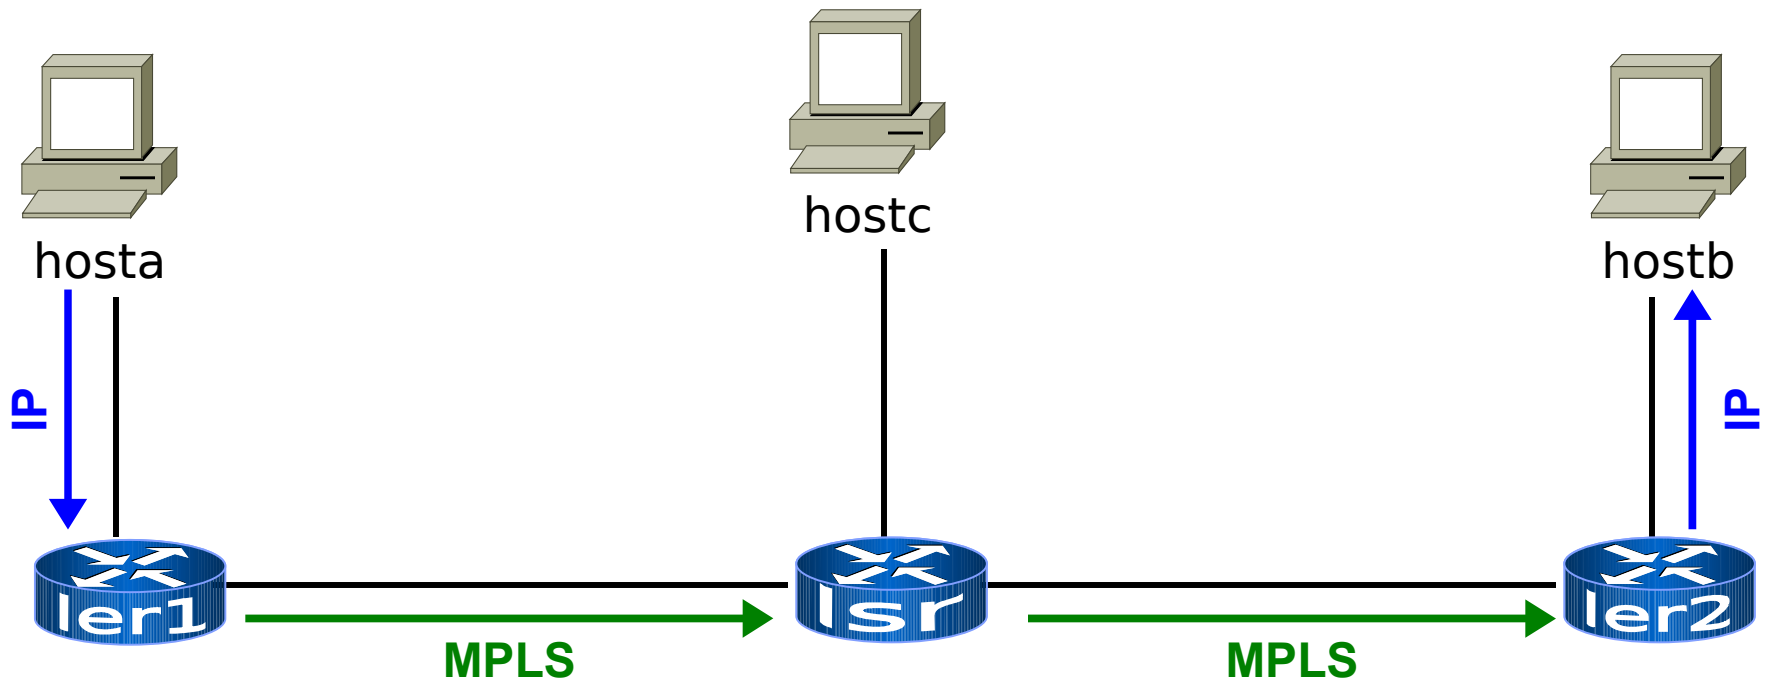
Network topology

- **hosta** is connected by an MPLS network (**1er1** – **1sr** – **1er2**) to **hostb**.
- **1er1**, **1sr**, and **1er2** must be configured in order to allow exchange of packets between **hosta** and **hostb** using MPLS.
- At the same time, **hosta** is connected to **hostc** via IP only. Therefore, IP routes have to be added to **hosta**, **hostc**, **1er1**, and **1sr** in order to support the exchange of conventional IP traffic.

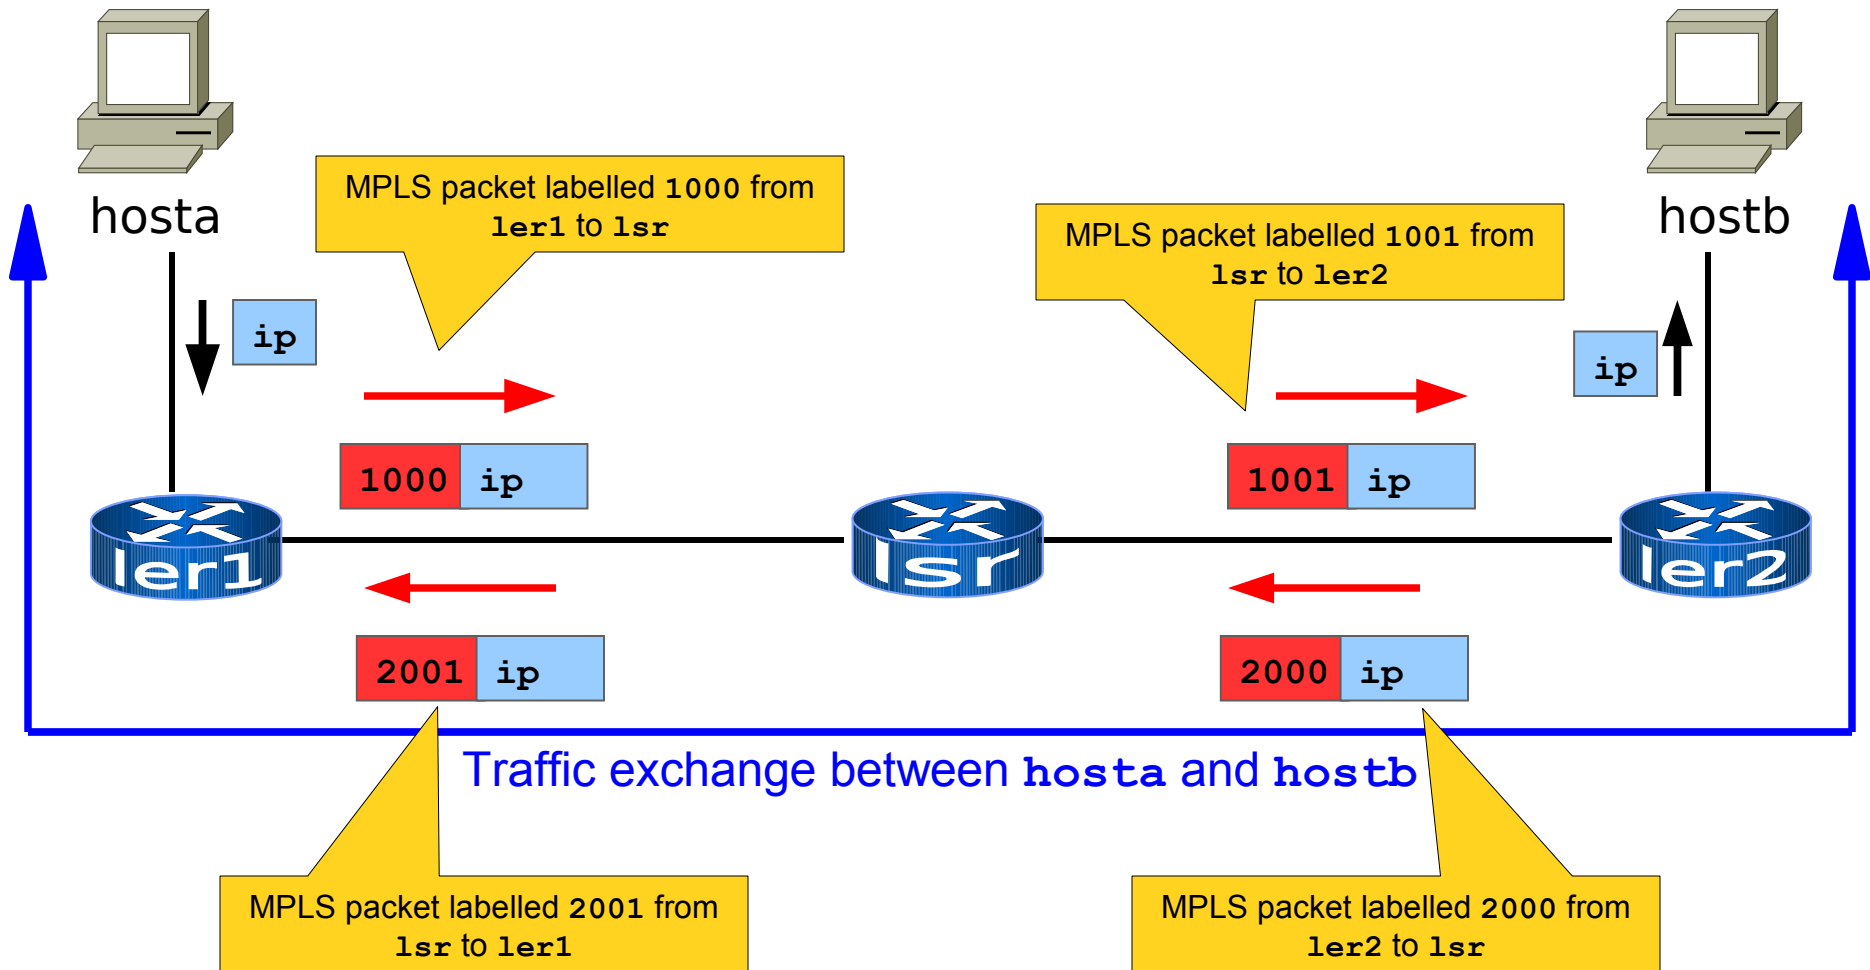
Traffic types



Traffic types



Traffic types



Basic configuration

hosta.startup

```
ifconfig eth1 192.168.0.1 netmask 255.255.255.0 broadcast 192.168.0.255 up
ifconfig dummy0 10.0.0.1/32
ip route add 10.0.0.2/32 via 192.168.0.2 src 10.0.0.1

#traffic to hostc
route add -net 172.16.2.0 netmask 255.255.255.0 gw 192.168.0.2 dev eth1
```

hostb.startup

```
ifconfig eth0 192.168.1.5 netmask 255.255.255.0 broadcast 192.168.1.255 up
ifconfig dummy0 10.0.0.2/32
ip route add 10.0.0.1/32 via 192.168.1.4 src 10.0.0.2
```

hostc.startup

```
ifconfig eth0 172.16.2.2 netmask 255.255.255.0 broadcast 172.16.2.255 up

#traffic to hosta
route add default gw 172.16.2.1
```

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlf add key 0 instructions \  
    push len 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add .0.0.2/32 via 172.16.0.3 mpls 0x2
```

new nhlf entry

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlf add key 0 instructions \  
    push gen 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add 0.2/32 via 172.16.0.3 mpls 0x2
```

sequential number identifying
the entry (0="new entry": a
number will be automatically
assigned)

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlfe add key 0 instructions \  
    push gen 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add 10.0.0.2/32 via 172.16.0.3 mpls 0x2
```

push a label of type "gen" and
value 1000...

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlfe add key 0 instructions \  
    push gen 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add 10.0.0.2/24 via 172.16.0.3 mpls 0x2
```

...and forward the packet to a
certain router

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlfe add key 0 instructions \  
    push gen 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add 10.0.0.2/24 via 172.16.0.3 mpls 0x2
```

...and forward the packet to a
certain router

Hey! But...
Isn't this plain
IP routing?

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlf add key 0 instructions \  
    push gen 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add 10.0.0.2/24 via 172.16.0.3 mpls 0x2
```

...and forward the packet to a
certain router

Hey! But...
Isn't this plain
IP routing?

Explanation
coming in
next slide...

MPLS configuration: 1er1

- Machines outside the MPLS network...
 - route by IP addresses
 - only know paths to ingress points to the MPLS network
- Routers inside the MPLS network...
 - route by label swapping
 - only know how to reach edge routers in the MPLS network...
 - ...and they achieve this knowledge by borrowing information from the IP routing tables built by an underlying IGP protocol (e.g., RIP, OSPF, etc.)
 - in this case, we simply use static routes
- This allows a complete separation of the routing inside and outside the MPLS network

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlfe add key 0 instructions \  
    push gen 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add 10.0.0.2/32 via 172.16.0.3 mpls 0x2
```

label binding: instruct the router to use the previously created nhlfe to forward the packet

MPLS configuration: 1er1

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

load MPLS
modules

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB *****
```

```
mpls nhlfe add key 0 instructions \  
    push gen 1000 nexthop eth2 ipv4 172.16.0.3
```

```
ip route add 10.0.0.2/32 via 172.16.0.3 mpls 0x2
```

this is the key returned by the
previous `mpls nhlfe add`
command

note: here we are defining a fec

MPLS configuration: 1er1

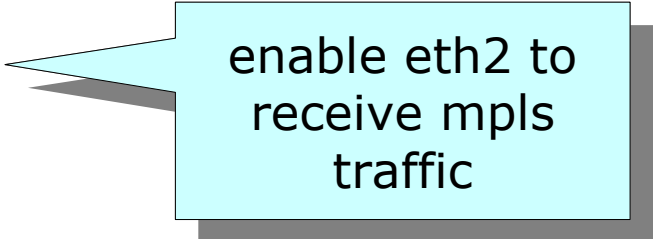
```
***** MPLS TRAFFIC FROM HOSTB TO HOSTA *****
```

```
mpls labelspace set dev eth2 labelspace 0
```

```
mpls ilm add label gen 2001 labelspace 0
```

```
mpls nhlf add key 0 instructions nexthop eth1 ipv4 192.168.0.1
```

```
mpls xc add ilm_label gen 2001 ilm_labelspace 0 nhlf_key 0x3
```



enable eth2 to
receive mpls
traffic

MPLS configuration: 1er1

```
***** MPLS TRAFFIC FROM HOSTB TO HOSTA *****
```

```
mpls labelspace set dev eth2 labelspace 0
```

```
mpls ilm add label gen 2001 labelspace 0
```

```
mpls nhlf add key 0 instructions nexthop eth2
```

```
mpls xc add ilm_label gen 2001 ilm_labelspace 0
```

put label 2001 in the ilm

required in order to recognize the incoming label and be able to put an entry in the xc list later on

MPLS configuration: 1er1

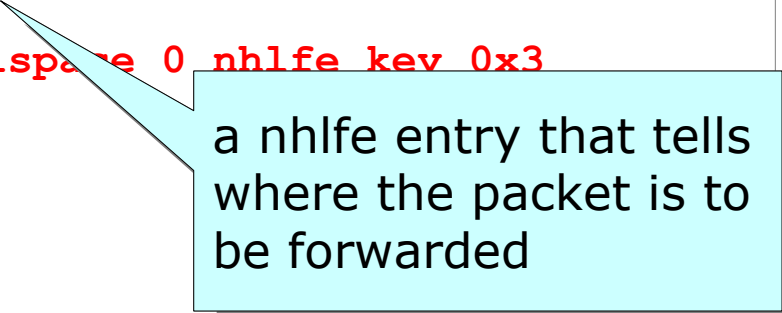
```
***** MPLS TRAFFIC FROM HOSTB TO HOSTA *****
```

```
mpls labelspace set dev eth2 labelspace 0
```

```
mpls ilm add label gen 2001 labelspace 0
```

```
mpls nhlf add key 0 instructions nexthop eth1 ipv4 192.168.0.1
```

```
mpls xc add ilm_label gen 2001 ilm_labelspace 0 nhlf key 0x3
```



a nhlf entry that tells where the packet is to be forwarded

MPLS configuration: 1er1

```
***** MPLS TRAFFIC FROM HOSTB TO HOSTA *****
```

```
mpls labelspace set dev eth2 labelspace 0
```

```
mpls ilm add label gen 2001 labelspace 0
```

```
mpls nhlfe add key 0 instructions nexthop eth1 ipv4 192.168.0.1
```

```
mpls xc add ilm_label gen 2001 ilm_labelspace 0 nhlfe_key 0x3
```

perform label "swapping":
upon receiving a packet with label 2001 we execute the nhlfe indexed by key 0x3 (returned by the last `mpls nhlfe add` command)

note: this instruction "consumes" (=pops) label 2001, which is all we need to do because packets directed to `hosta` are exiting the MPLS network

MPLS configuration: `1er2`

- Very similar to `1er1`'s

MPLS configuration: 1sr

```
modprobe mpls4  
modprobe mplsbr  
modprobe mpls_tunnel
```

```
***** MPLS TRAFFIC FROM HOSTA TO HOSTB
```

```
mpls labelspace set dev eth2 labelspace 0  
mpls ilm add label gen 1000 labelspace 0  
mpls nhlfe add key 0 instructions \  
    push gen 1001 nexthop eth3 ipv4 172.16.1.4  
mpls xc add ilm_label gen 1000 ilm_labelspace 0 nhlfe_key 0x2
```

swap incoming label
1000 with 1001 and
forward on to 1er2

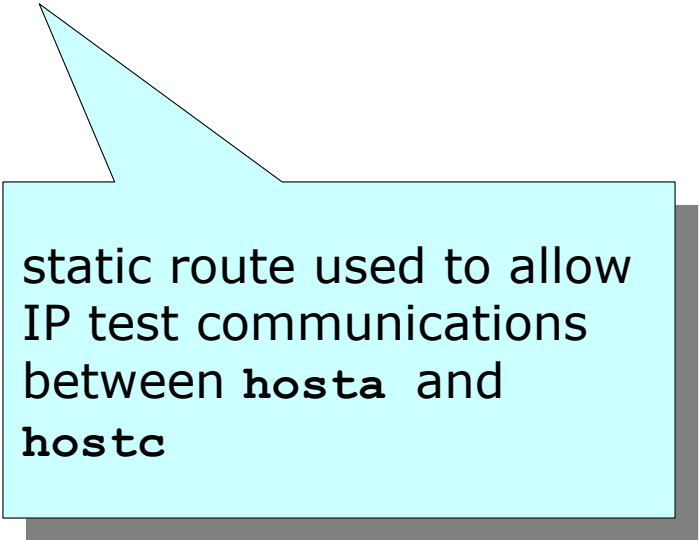
```
***** MPLS TRAFFIC FROM HOSTB TO HOSTA
```

```
mpls labelspace set dev eth3 labelspace 0  
mpls ilm add label gen 2000 labelspace 0  
mpls nhlfe add key 0 instructions \  
    push gen 2001 nexthop eth2 ipv4 172.16.0.2  
mpls xc add ilm_label gen 2000 ilm_labelspace 0 nhlfe_key 0x3
```

swap incoming label
2000 with 2001 and
forward on to 1er1

MPLS configuration: 1sr

```
route add -net 192.168.0.0 netmask 255.255.255.0 gw  
172.16.0.2 dev eth2
```



static route used to allow
IP test communications
between **hosta** and
hostc

Starting the lab

host machine

```
user@localhost:~$ cd netkit-lab_mpls  
user@localhost:~/netkit-lab_mpls$ lstart
```

Commands for MPLS traffic analysis

hosta

```
hosta:~# ping 10.0.0.2
```

lsr

```
lsr:~# tcpdump -i eth3
```

ler2

```
ler2:~# tcpdump -i eth0
```

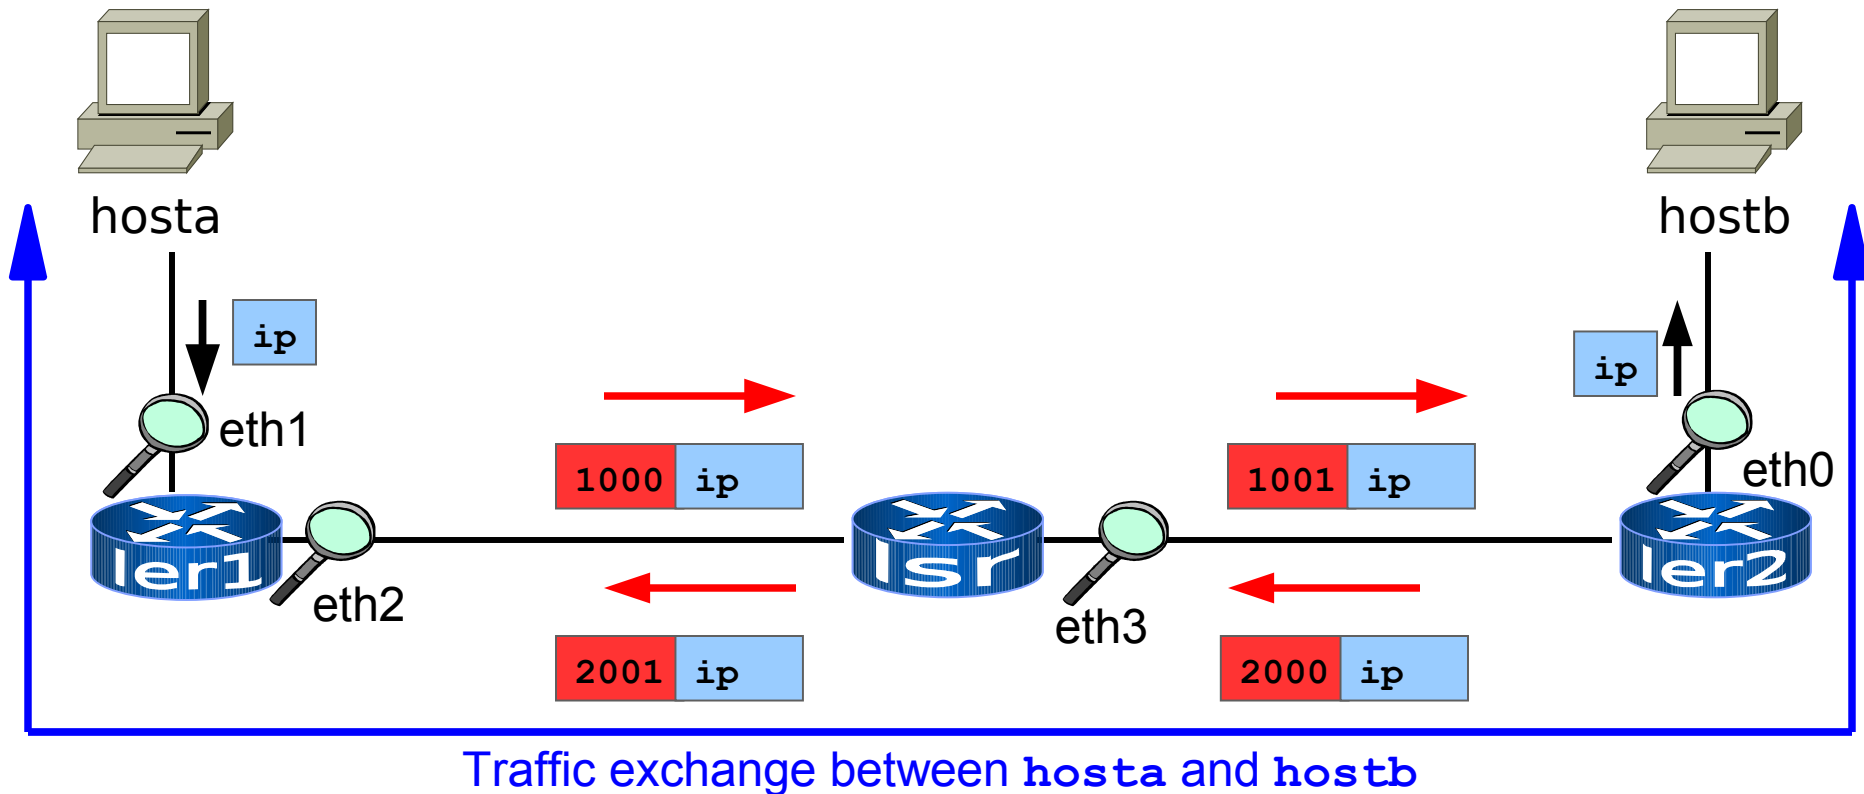
ler1

```
ler1:~# tcpdump -i eth1
```

ler1

```
ler1:~# tcpdump -i eth2
```

MPLS Traffic analysis



MPLS Traffic analysis

- Some checkpoints have been identified along the path from **hosta** to **hostb**
 - **ler1**: both interfaces, to observe traffic before and after the insertion of the MPLS header
 - **lsr**: interface **eth3**, to observe label switching
 - **ler2**: interface **eth0**, where the MPLS header is removed
- Similar checkpoints can be considered for the traffic from **hostb** to **hosta**

MPLS Traffic analysis

- **ler1** applies NHLFE 0x2 (=“push label 1000”) to traffic directed to 10.0.0.2 and forwards it to **lsr**

▼ **ler1**

```
ler1:~# ip route show
10.0.0.2 via 172.16.0.3 dev eth2 mpls 0x2
172.16.2.0/24 via 172.16.0.3 dev eth2
172.16.0.0/24 dev eth2 proto kernel scope link src 172.16.0.2
192.168.0.0/24 dev eth1 proto kernel scope link src 192.168.0.2
ler1:~# mpls nhlfe show
NHLFE entry key 0x00000003 mtu 1500 propagate_ttl
      set eth1 ipv4 192.168.0.1 (168 bytes, 2 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
      push gen 1000 set eth2 ipv4 172.16.0.3 (168 bytes, 2 pkts)
ler1:~# █
```


MPLS Traffic analysis

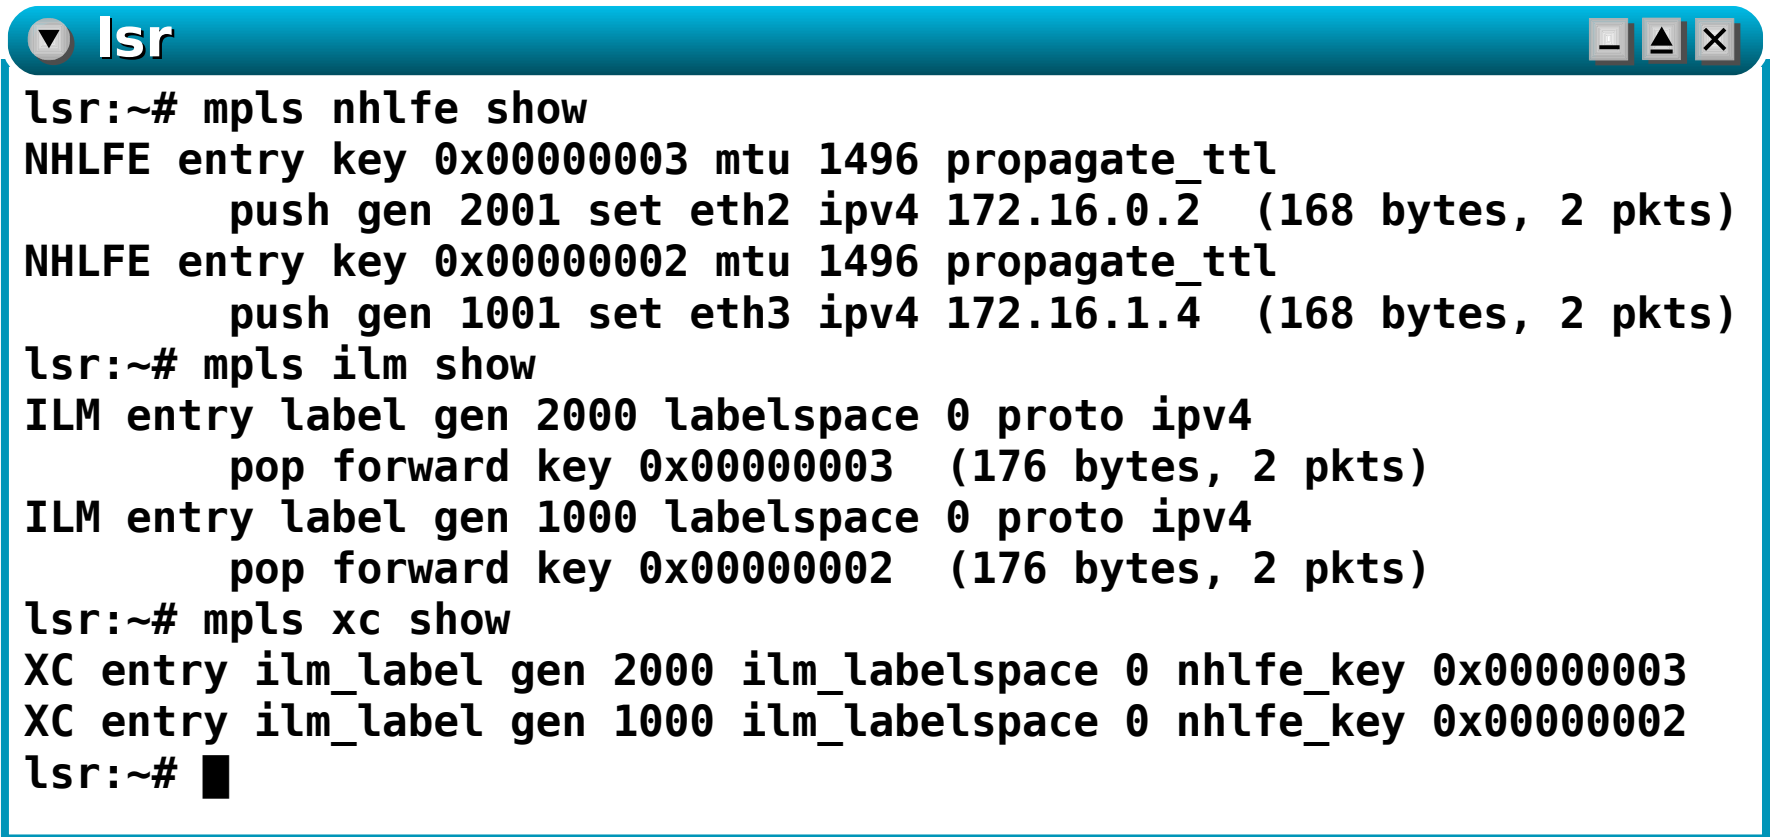
- `lsr`'s IP (correctly) knows nothing about `10.0.0.0s`, but...

lsr

```
lsr:~# ip route show
172.16.2.0/24 dev eth0 proto kernel scope link src 172.16.2.1
172.16.0.0/24 dev eth2 proto kernel scope link src 172.16.0.3
192.168.0.0/24 via 172.16.0.2 dev eth2
172.16.1.0/24 dev eth3 proto kernel scope link src 172.16.1.3
lsr:~# █
```

MPLS Traffic analysis

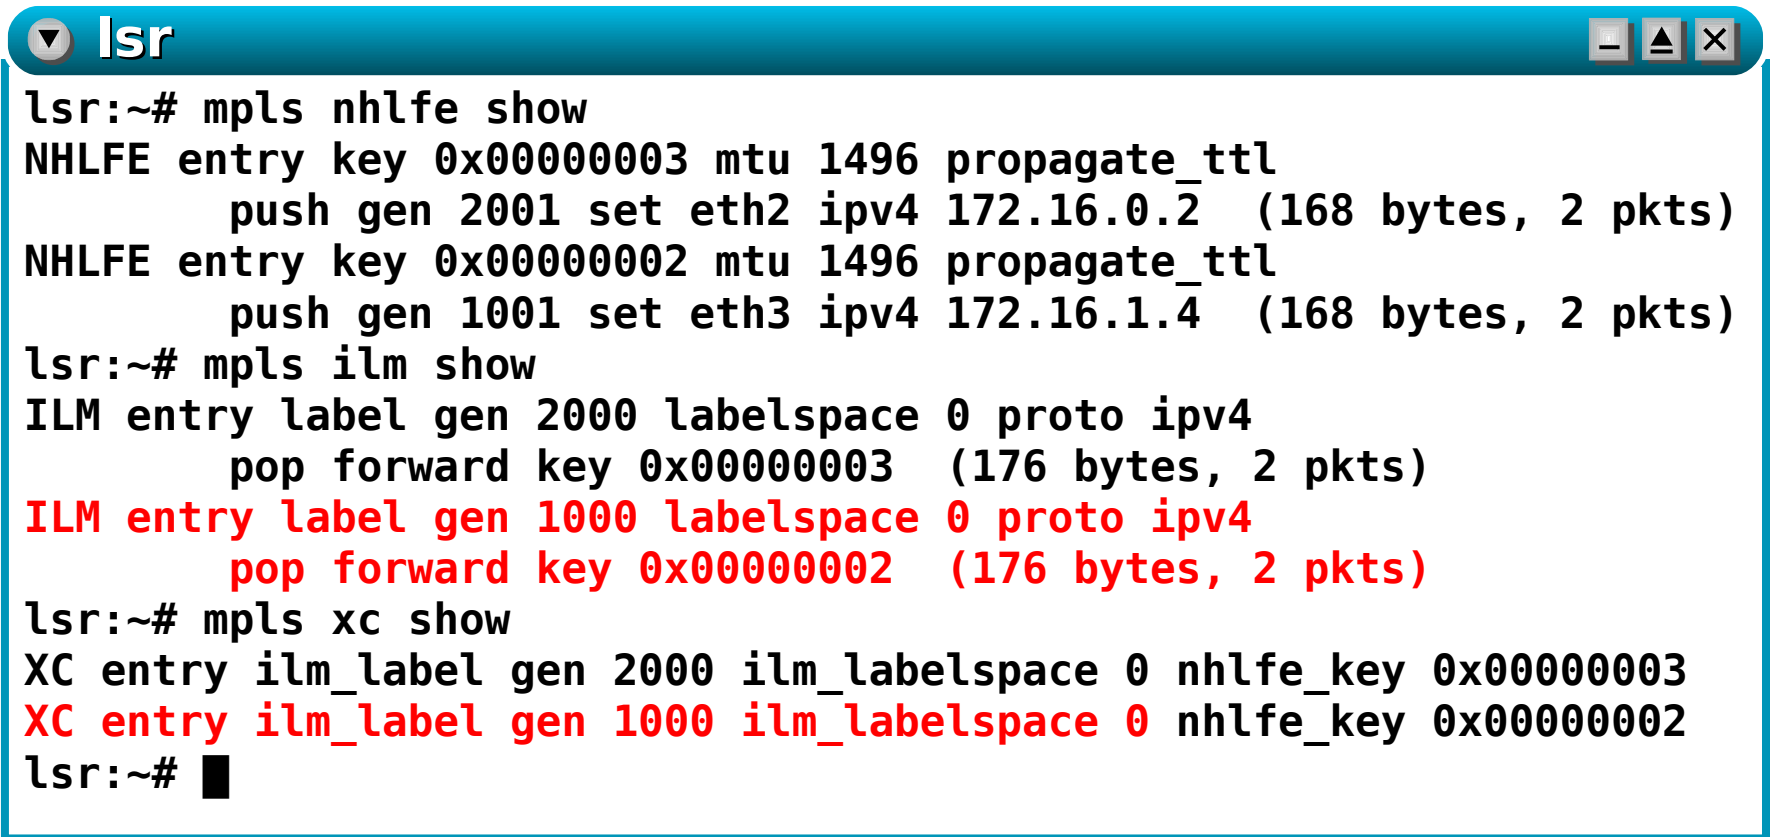
- ...**lsr**'s MPLS knows how to forward labeled packets!



```
lsr:~# mpls nhlfe show
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 2001 set eth2 ipv4 172.16.0.2 (168 bytes, 2 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
    push gen 1001 set eth3 ipv4 172.16.1.4 (168 bytes, 2 pkts)
lsr:~# mpls ilm show
ILM entry label gen 2000 labelspace 0 proto ipv4
    pop forward key 0x00000003 (176 bytes, 2 pkts)
ILM entry label gen 1000 labelspace 0 proto ipv4
    pop forward key 0x00000002 (176 bytes, 2 pkts)
lsr:~# mpls xc show
XC entry ilm_label gen 2000 ilm_labelspace 0 nhlfe_key 0x00000003
XC entry ilm_label gen 1000 ilm_labelspace 0 nhlfe_key 0x00000002
lsr:~# █
```

MPLS Traffic analysis

- incoming packets with label 1000 have their label recognized and popped...



```
lsr:~# mpls nhlfe show
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 2001 set eth2 ipv4 172.16.0.2 (168 bytes, 2 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
    push gen 1001 set eth3 ipv4 172.16.1.4 (168 bytes, 2 pkts)
lsr:~# mpls ilm show
ILM entry label gen 2000 labelspace 0 proto ipv4
    pop forward key 0x00000003 (176 bytes, 2 pkts)
ILM entry label gen 1000 labelspace 0 proto ipv4
    pop forward key 0x00000002 (176 bytes, 2 pkts)
lsr:~# mpls xc show
XC entry ilm_label gen 2000 ilm_labelspace 0 nhlfe_key 0x00000003
XC entry ilm_label gen 1000 ilm_labelspace 0 nhlfe_key 0x00000002
lsr:~# █
```

MPLS Traffic analysis

- ...and are forwarded to **1er2** after swapping the label with **1001**

```
lslr
lslr:~# mpls nhlfe show
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 2001 set eth2 ipv4 172.16.0.2 (168 bytes, 2 pkts)
NHLFE entry key 0x00000002 mtu 1496 propagate_ttl
    push gen 1001 set eth3 ipv4 172.16.1.4 (168 bytes, 2 pkts)
lslr:~# mpls ilm show
ILM entry label gen 2000 labelspace 0 proto ipv4
    pop forward key 0x00000003 (176 bytes, 2 pkts)
ILM entry label gen 1000 labelspace 0 proto ipv4
    pop forward key 0x00000002 (176 bytes, 2 pkts)
lslr:~# mpls xc show
XC entry ilm_label gen 2000 ilm_labelspace 0 nhlfe_key 0x00000003
XC entry ilm_label gen 1000 ilm_labelspace 0 nhlfe_key 0x00000002
lslr:~# █
```

MPLS Traffic analysis

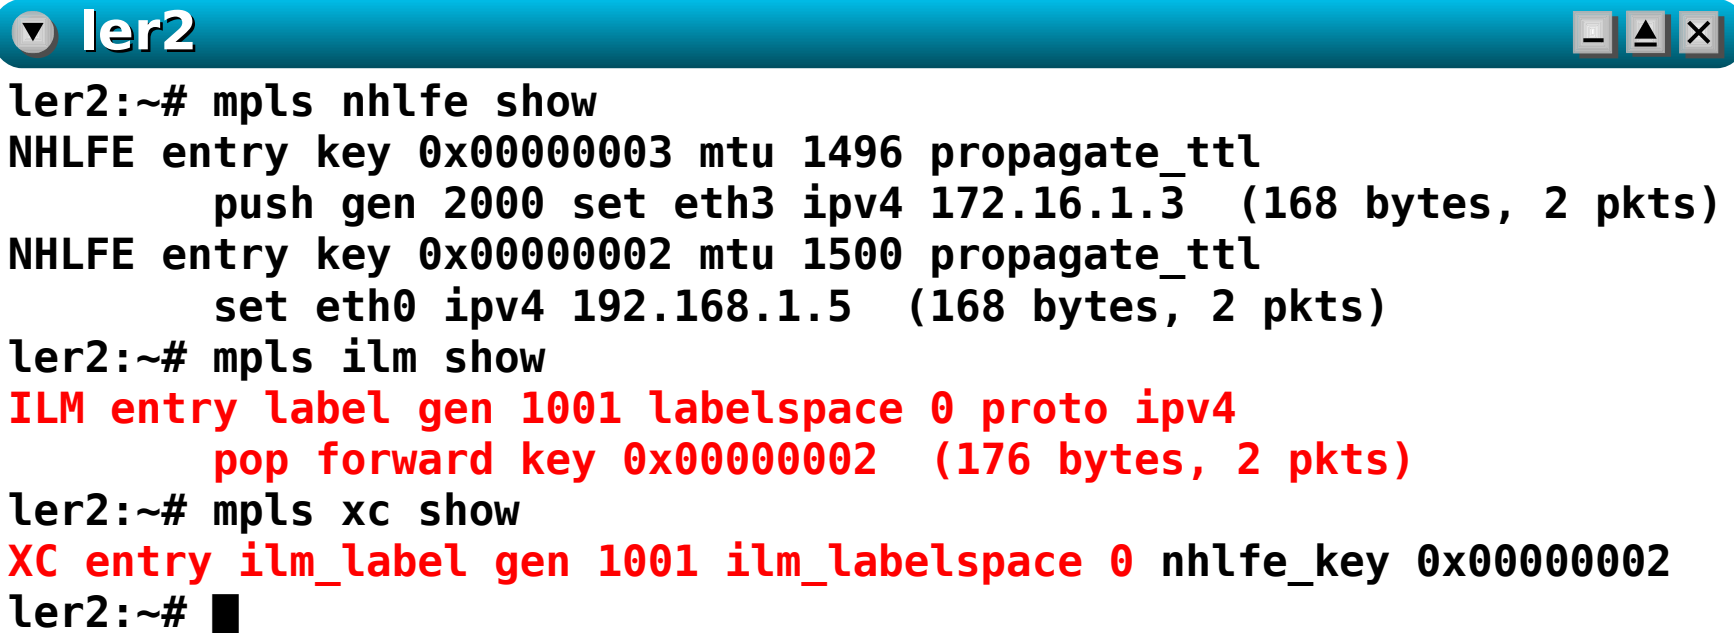
- also `ler2`'s IP (correctly) knows nothing about `10.0.0.2`, but...

ler2

```
ler2:~# ip route show
10.0.0.1 via 172.16.1.3 dev eth3 mpls 0x3
192.168.1.0/24 dev eth0 proto kernel scope link src 192.168.1.4
172.16.1.0/24 dev eth3 proto kernel scope link src 172.16.1.4
ler2:~# █
```

MPLS Traffic analysis

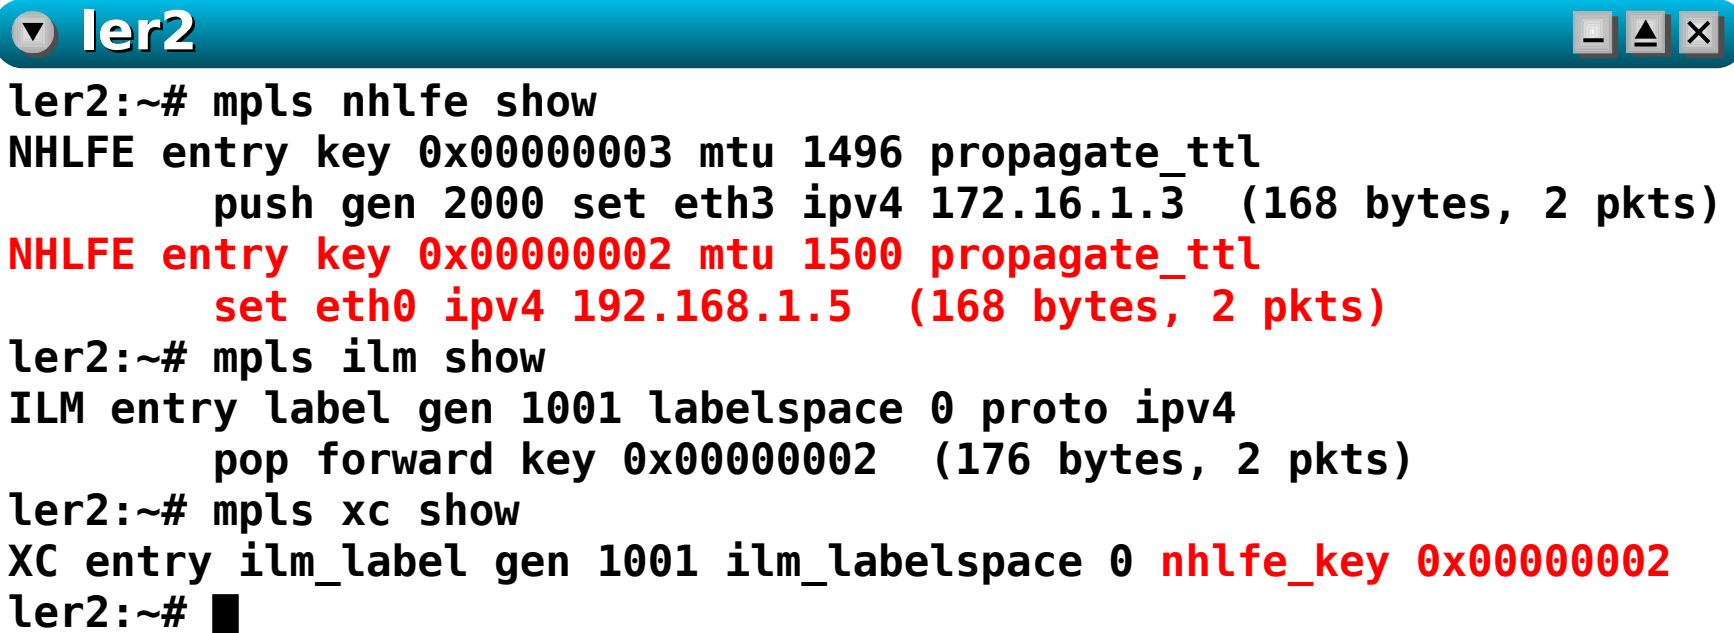
- ...incoming packets with label 1001 have their label recognized and popped...



```
ler2:~# mpls nhlfe show
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 2000 set eth3 ipv4 172.16.1.3 (168 bytes, 2 pkts)
NHLFE entry key 0x00000002 mtu 1500 propagate_ttl
    set eth0 ipv4 192.168.1.5 (168 bytes, 2 pkts)
ler2:~# mpls ilm show
ILM entry label gen 1001 labelspace 0 proto ipv4
    pop forward key 0x00000002 (176 bytes, 2 pkts)
ler2:~# mpls xc show
XC entry ilm_label gen 1001 ilm_labelspace 0 nhlfe_key 0x00000002
ler2:~# █
```

MPLS Traffic analysis

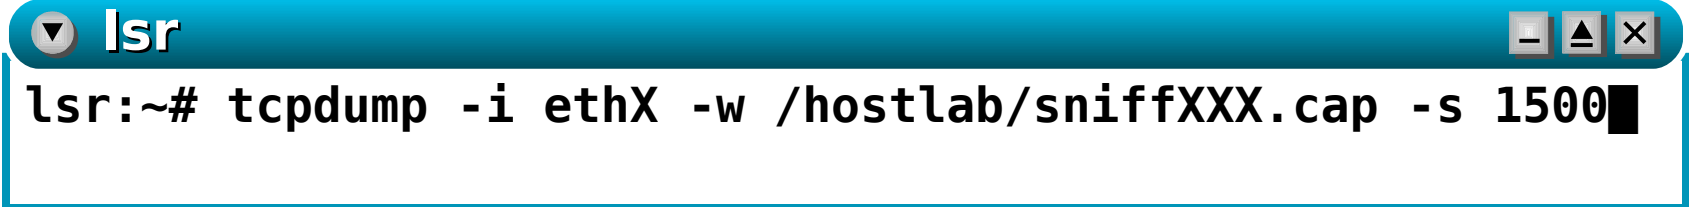
- ...and are forwarded to hostb



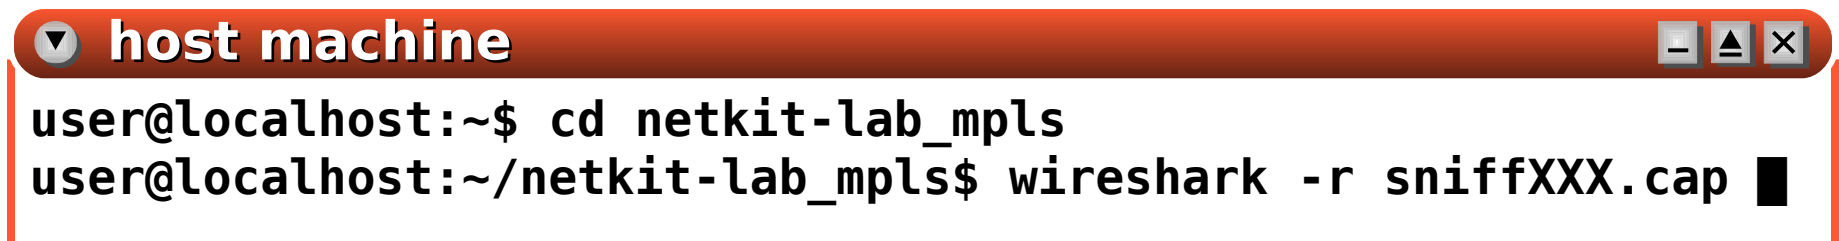
```
ler2:~# mpls nhlfe show
NHLFE entry key 0x00000003 mtu 1496 propagate_ttl
    push gen 2000 set eth3 ipv4 172.16.1.3 (168 bytes, 2 pkts)
NHLFE entry key 0x00000002 mtu 1500 propagate_ttl
    set eth0 ipv4 192.168.1.5 (168 bytes, 2 pkts)
ler2:~# mpls ilm show
ILM entry label gen 1001 labelspace 0 proto ipv4
    pop forward key 0x00000002 (176 bytes, 2 pkts)
ler2:~# mpls xc show
XC entry ilm_label gen 1001 ilm_labelspace 0 nhlfe_key 0x00000002
ler2:~# █
```

MPLS Traffic analysis

- We sniff packets using tcpdump and examine the dumps on the host using wireshark

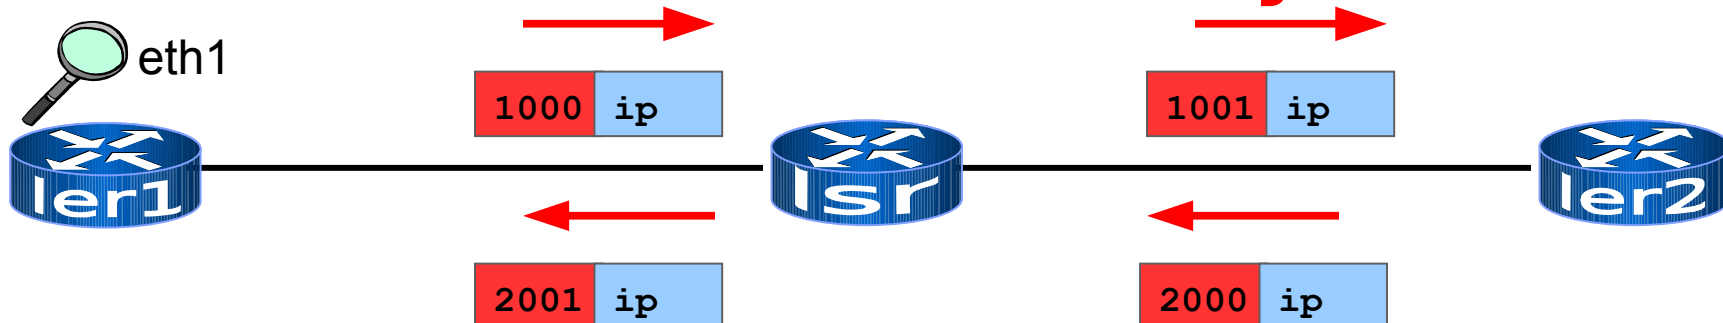


```
lsr:~# tcpdump -i ethX -w /hostlab/sniffXXX.cap -s 1500
```



```
user@localhost:~$ cd netkit-lab_mpls
user@localhost:~/netkit-lab_mpls$ wireshark -r sniffXXX.cap
```

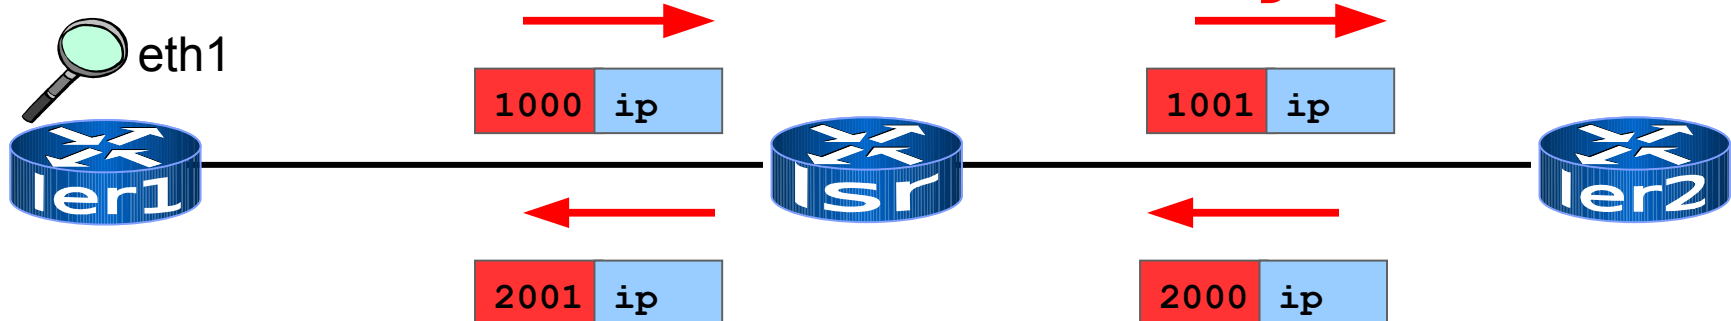

MPLS Traffic analysis



No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000540	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.010003	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010808	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	2.020014	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
6	2.020541	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
7	3.030045	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
8	3.030628	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
9	4.040040	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	4.040623	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	5.050014	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	5.050569	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

Frame 1 (98 bytes on wire, 98 bytes captured)
Ethernet II, Src: 4a:b6:3f:db:98:62 (4a:b6:3f:db:98:62), Dst: 5e:93:ba:fa:7c:1c (5e:93:ba:fa:7c:1c)
Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
Internet Control Message Protocol

MPLS Traffic analysis

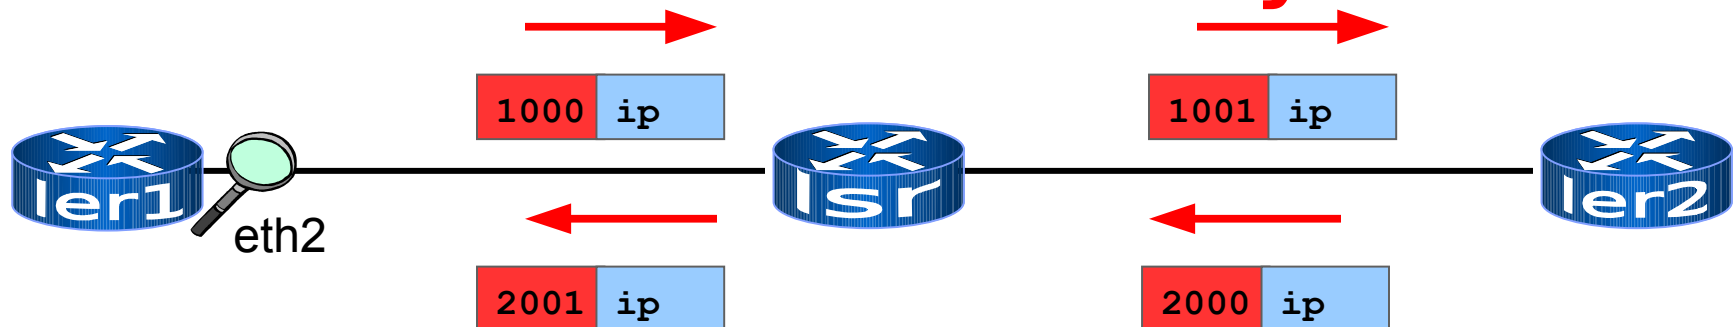


No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000540	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.010003	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010808	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	2.020014	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
6	2.020541	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
7	3.030045	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
8	3.030628	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
9	4.040040	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	4.040623	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	5.050014	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	5.050569	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

plain ICMP packets

Frame 1 (98 bytes on wire, 98 bytes captured)
Ethernet II, Src: 4a:b6:3f:db:98:62 (4a:b6:3f:db:98:62), Dst: 5e:93:ba:fa:7c:1c (5e:93:ba:fa:7c:1c)
Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
Internet Control Message Protocol

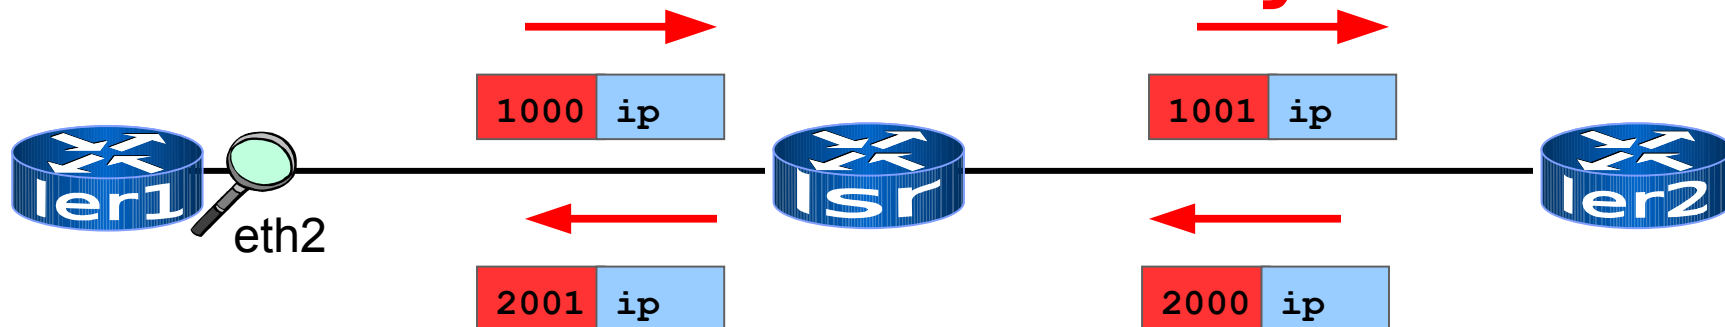
MPLS Traffic analysis



No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000440	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.009669	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010171	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	2.009923	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
6	2.010422	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
7	2.879181	86:59:f1:c8:89:f8	6a:a8:77:72:7a:d8	ARP	Who has 172.16.0.2? Te
8	2.879205	6a:a8:77:72:7a:d8	86:59:f1:c8:89:f8	ARP	172.16.0.2 is at 6a:a8:
9	3.019863	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	3.020483	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	4.029677	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	4.030141	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

▶ Frame 1 (102 bytes on wire, 102 bytes captured)
▶ Ethernet II, Src: 6a:a8:77:72:7a:d8 (6a:a8:77:72:7a:d8), Dst: 86:59:f1:c8:89:f8 (86:59:f1:c8:89:f8)
▼ MultiProtocol Label Switching Header, Label: 1000, Exp: 0, S: 1, TTL: 63
MPLS Label: 1000
MPLS Experimental Bits: 0
MPLS Bottom Of Label Stack: 1
MPLS TTL: 63
▶ Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
▶ Internet Control Message Protocol

MPLS Traffic analysis



No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000440	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.009669	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010171	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	2.009923	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
6	2.010425	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
7	3.009669	10.0.0.1	10.0.0.2	ARP	Who has 172.16.0.2? Te
8	3.010171	10.0.0.2	10.0.0.1	ARP	172.16.0.2 is at 6a:a8:
9	4.009669	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	4.010171	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	4.029669	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	4.030141	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

ICMP packets...

Frame 1 (102 bytes on wire, 102 bytes captured)

Ethernet II, Src: 6a:a8:77:72:7a:d8 (6a:a8:77:72:7a:d8), Dst: 86:59:f1:c8:89:f8 (86:59:f1:c8:89:f8)

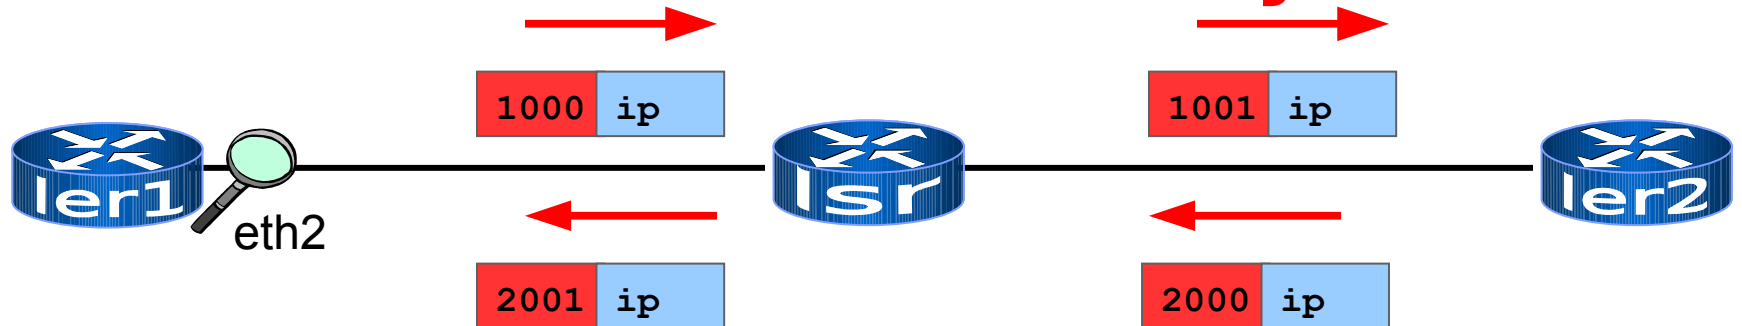
MultiProtocol Label Switching Header, Label: 1000, Exp: 0, S: 1, TTL: 63

- MPLS Label: 1000
- MPLS Experimental Bits: 0
- MPLS Bottom Of Label Stack: 1
- MPLS TTL: 63

Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)

Internet Control Message Protocol

MPLS Traffic analysis



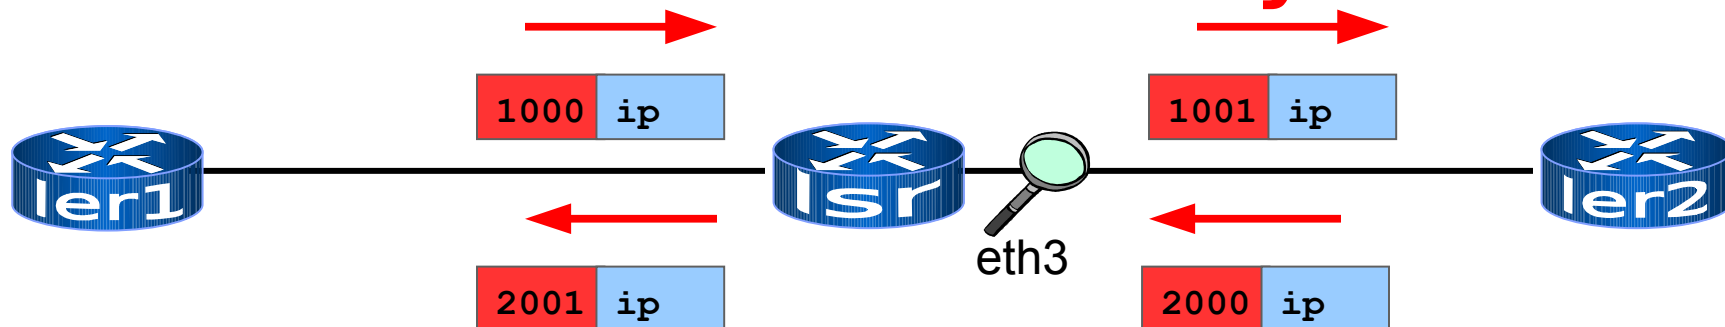
No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000440	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.009669	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010171	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	2.009923	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
6	2.010425	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
7	3.009669	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
8	3.010171	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
9	4.009669	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	4.010171	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	4.020171	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	4.030141	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

Frame 1 (102 bytes on wire, 102 bytes captured)
Ethernet II, Src: 6a:a8:77:72:7a:d8 (6a:a8:77:72:7a:d8), Dst: 86:59:f1:c8:89:f8 (86:59:f1:c8:89:f8)
MultiProtocol Label Switching Header, Label: 1000, Exp: 0, S: 1, TTL: 63
MPLS Label: 1000
MPLS Experimental Bits: 0
MPLS Bottom Of Label Stack: 1
MPLS TTL: 63
Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
Internet Control Message Protocol

ICMP
packets...

...encapsulated
in MPLS

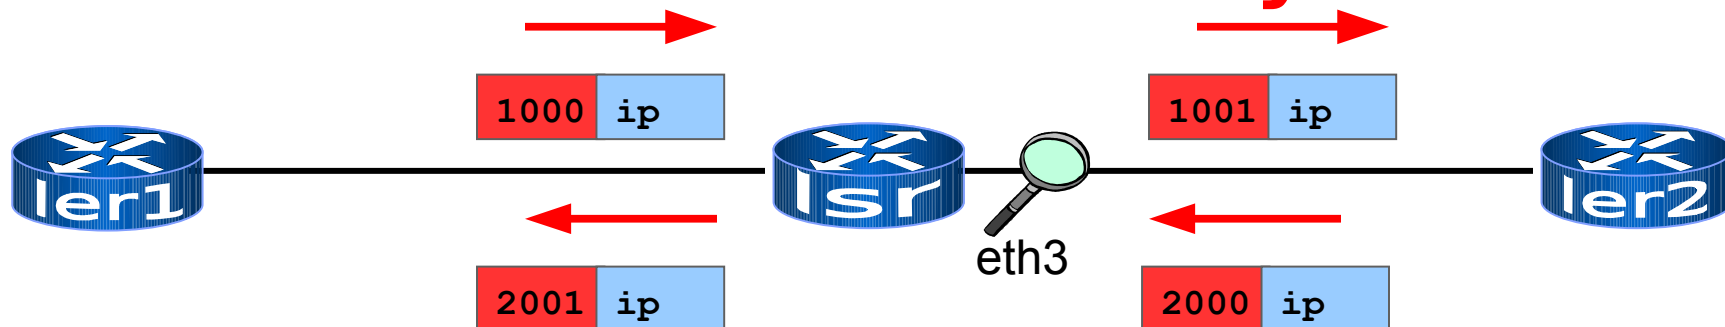
MPLS Traffic analysis



No.	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000313	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.009854	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010204	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	1.160649	aa:83:be:20:c1:bd	d2:8b:a4:86:7e:fd	ARP	Who has 172.16.1.4? Te
6	1.160859	d2:8b:a4:86:7e:fd	aa:83:be:20:c1:bd	ARP	172.16.1.4 is at d2:8b:
7	2.020003	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
8	2.020365	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
9	3.029933	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	3.030262	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	4.039937	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	4.040333	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

▸ Frame 1 (102 bytes on wire, 102 bytes captured)
▸ Ethernet II, Src: aa:83:be:20:c1:bd (aa:83:be:20:c1:bd), Dst: d2:8b:a4:86:7e:fd (d2:8b:a4:86:7e:fd)
▾ MultiProtocol Label Switching Header, Label: 1001, Exp: 0, S: 1, TTL: 62
MPLS Label: 1001
MPLS Experimental Bits: 0
MPLS Bottom Of Label Stack: 1
MPLS TTL: 62
▸ Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
▸ Internet Control Message Protocol

MPLS Traffic analysis



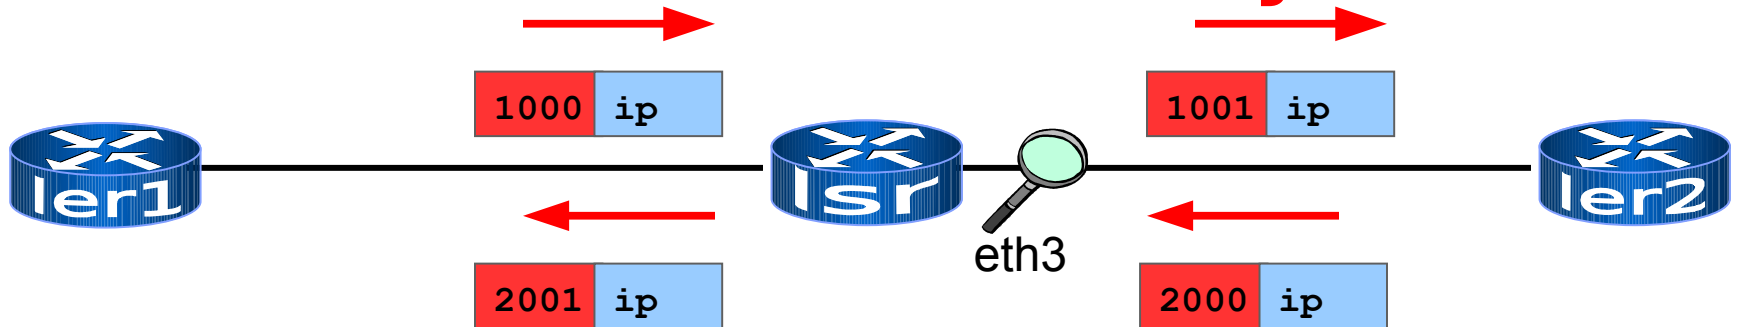
No.	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000313	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.009854	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010204	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	1.160649	aa:83:be:20:c1:bd	d2:8b:7e:fd	ARP	Who has 172.16.1.4? Te
6	1.160649	aa:83:be:20:c1:bd	d2:8b:7e:fd	ARP	172.16.1.4 is at d2:8b:7e:fd
7	1.160649	10.0.0.2	10.0.0.1	ICMP	Echo (ping) request
8	1.160649	10.0.0.1	10.0.0.2	ICMP	Echo (ping) reply
9	1.160649	10.0.0.2	10.0.0.1	ICMP	Echo (ping) request
10	1.160649	10.0.0.1	10.0.0.2	ICMP	Echo (ping) reply
11	4.039937	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	4.040333	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

ICMP packets...

...encapsulated in MPLS...

Frame 1 (102 bytes on wire, 102 bytes captured)
Ethernet II, Src: aa:83:be:20:c1:bd (aa:83:be:20:c1:bd), Dst: d2:8b:7e:fd (d2:8b:7e:fd)
MultiProtocol Label Switching Header, Label: 1001, Exp: 0, S: 1, TTL: 62
MPLS Label: 1001
MPLS Experimental Bits: 0
MPLS Bottom Of Label Stack: 1
MPLS TTL: 62
Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
Internet Control Message Protocol

MPLS Traffic analysis



No.	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000313	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.009854	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010204	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	1.160649	aa:83:be:20:c1:bd	d2:8b:a6:7e:fd	ARP	Who has 172.16.1.4? Te
6	1.160649	aa:83:be:20:c1:bd	d2:8b:a6:7e:fd	ARP	172.16.1.4 is at d2:8b:
7	1.160649	10.0.0.2	10.0.0.1	ICMP	Echo (ping) request
8	1.160649	10.0.0.1	10.0.0.2	ICMP	Echo (ping) reply
9	1.160649	10.0.0.2	10.0.0.1	ICMP	Echo (ping) request
10	1.160649	10.0.0.1	10.0.0.2	ICMP	Echo (ping) reply
11	4.039937	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	4.040333	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

ICMP packets...

...encapsulated in MPLS...

MultiProtocol Label Switching Header, Label: 1001, Exp: 0, S: 1, TTL: 62

MPLS Label: 1001

MPLS Experimental Bits: 0

MPLS Bottom Of Label Stack: 1

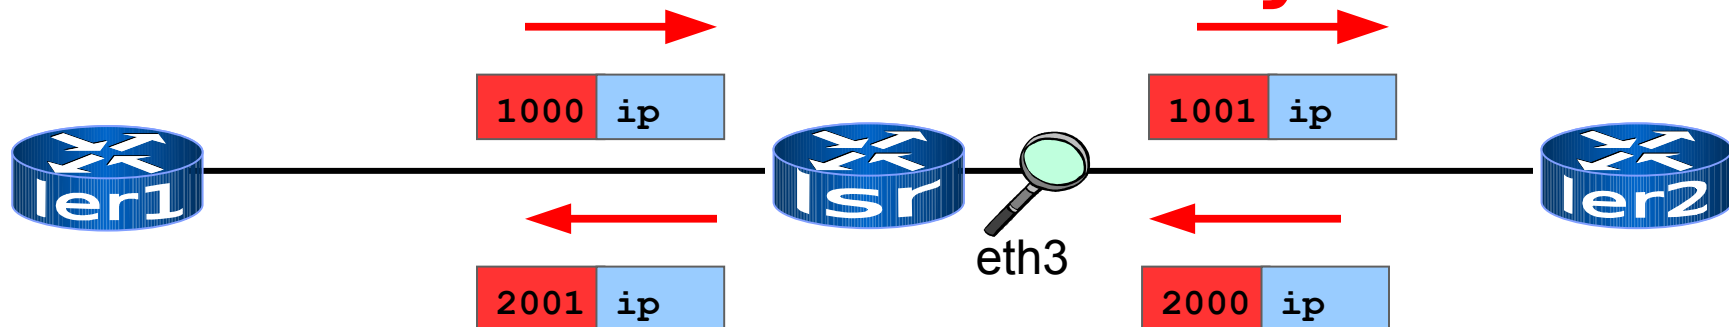
MPLS TTL: 62

Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2

Internet Control Message Protocol

...with a swapped label!

MPLS Traffic analysis

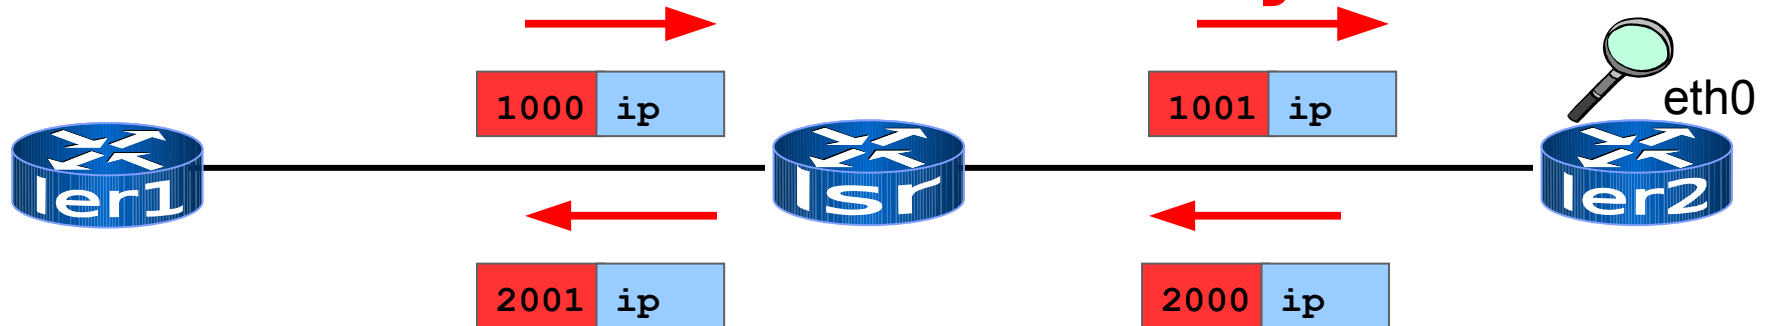


No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000313	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.009854	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010204	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	1.160649	aa:83:be:20:c1:bd	d2:8b:a4:86:7e:fd	ARP	Who has 172.16.1.4? Te
6		aa:83:be:20:c1:bd		ARP	172.16.1.4 is at d2:8b:
		.0.0.2	.0.0.1	ICMP	Echo (ping) request
		.0.0.1	.0.0.2	ICMP	Echo (ping) reply
		.0.0.2	.0.0.1	ICMP	Echo (ping) request
		.0.0.1	.0.0.2	ICMP	Echo (ping) reply
		.0.0.2	.0.0.1	ICMP	Echo (ping) request
		.0.0.1	.0.0.2	ICMP	Echo (ping) reply

echo replies travel with a different label

Frame 2 (102 bytes on wire, 8 bytes captured)
Ethernet II, Src: d2:8b:a4:86:7e:fd (d2:8b:a4:86:7e:fd), Dst: aa:83:be:20:c1:bd (aa:83:be:20:c1:bd)
MultiProtocol Label Switching Header, Label: 2000, Exp: 0, S: 1, TTL: 63
MPLS Label: 2000
MPLS Experimental Bits: 0
MPLS Bottom Of Label Stack: 1
MPLS TTL: 63
Internet Protocol, Src: 10.0.0.2 (10.0.0.2), Dst: 10.0.0.1 (10.0.0.1)
Internet Control Message Protocol

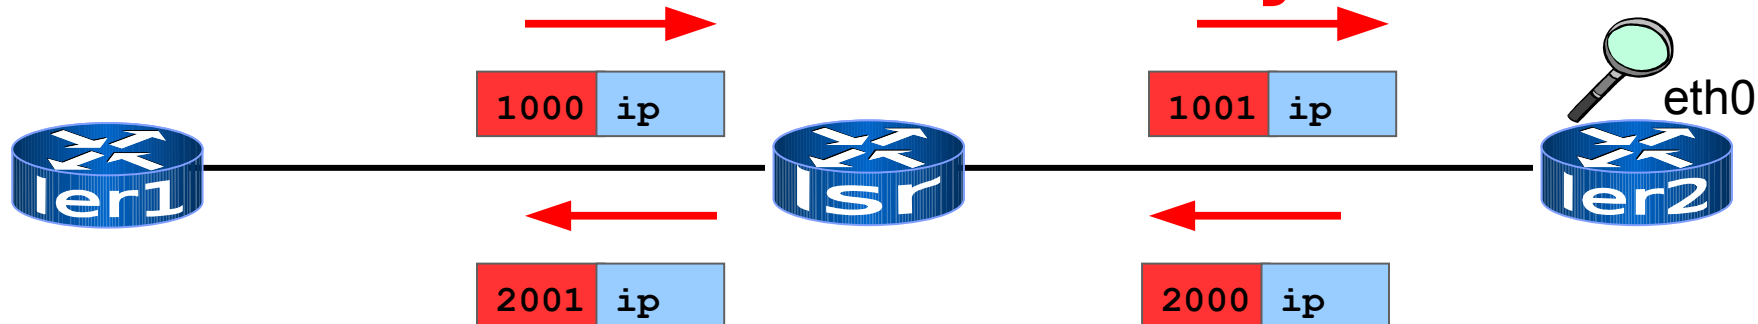
MPLS Traffic analysis



No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000121	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.010061	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010313	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	2.020049	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
6	2.020245	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
7	3.030085	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
8	3.030287	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
9	4.040128	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	4.040307	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	5.049959	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	5.050155	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

Frame 1 (98 bytes on wire, 98 bytes captured)
Ethernet II, Src: ae:48:62:f5:63:c6 (ae:48:62:f5:63:c6), Dst: de:e5:92:ad:6f:e4 (d
Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
Internet Control Message Protocol

MPLS Traffic analysis

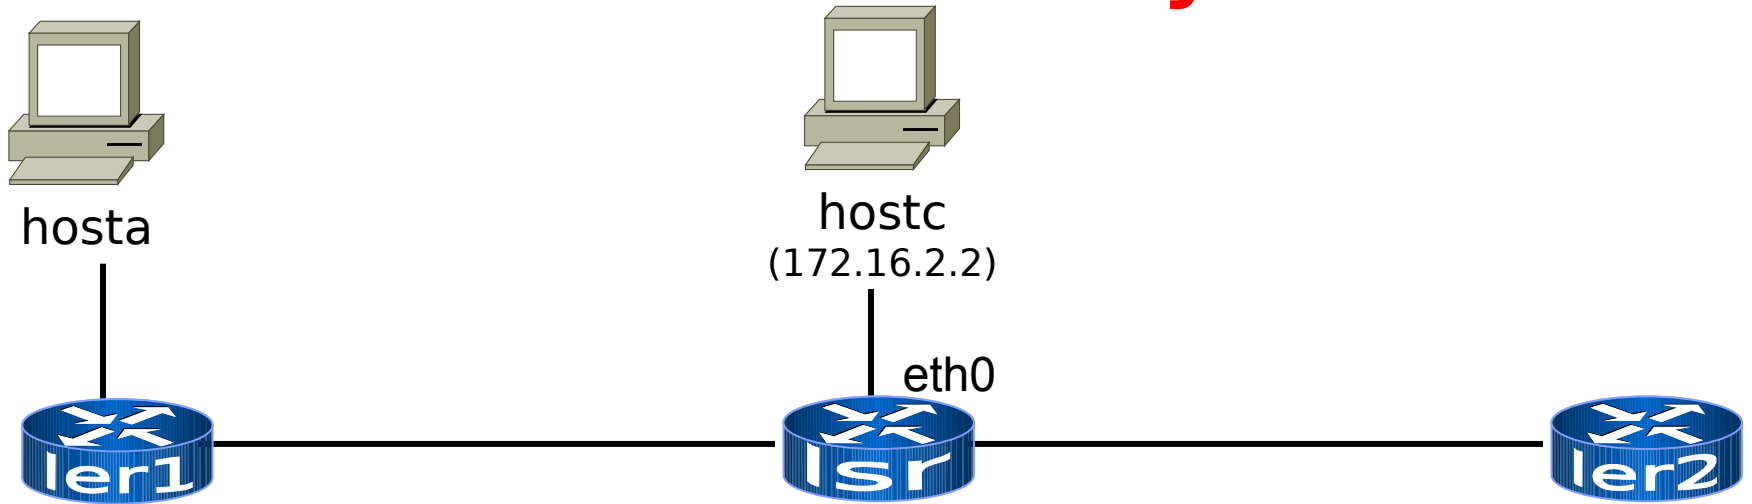


No. .	Time	Source	Destination	Protocol	Info
1	0.000000	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
2	0.000121	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
3	1.010061	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
4	1.010313	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
5	2.020049	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
6	2.020245	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
7	3.030085	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
8	3.030287	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
9	4.040128	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
10	4.040307	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply
11	5.049959	10.0.0.1	10.0.0.2	ICMP	Echo (ping) request
12	5.050155	10.0.0.2	10.0.0.1	ICMP	Echo (ping) reply

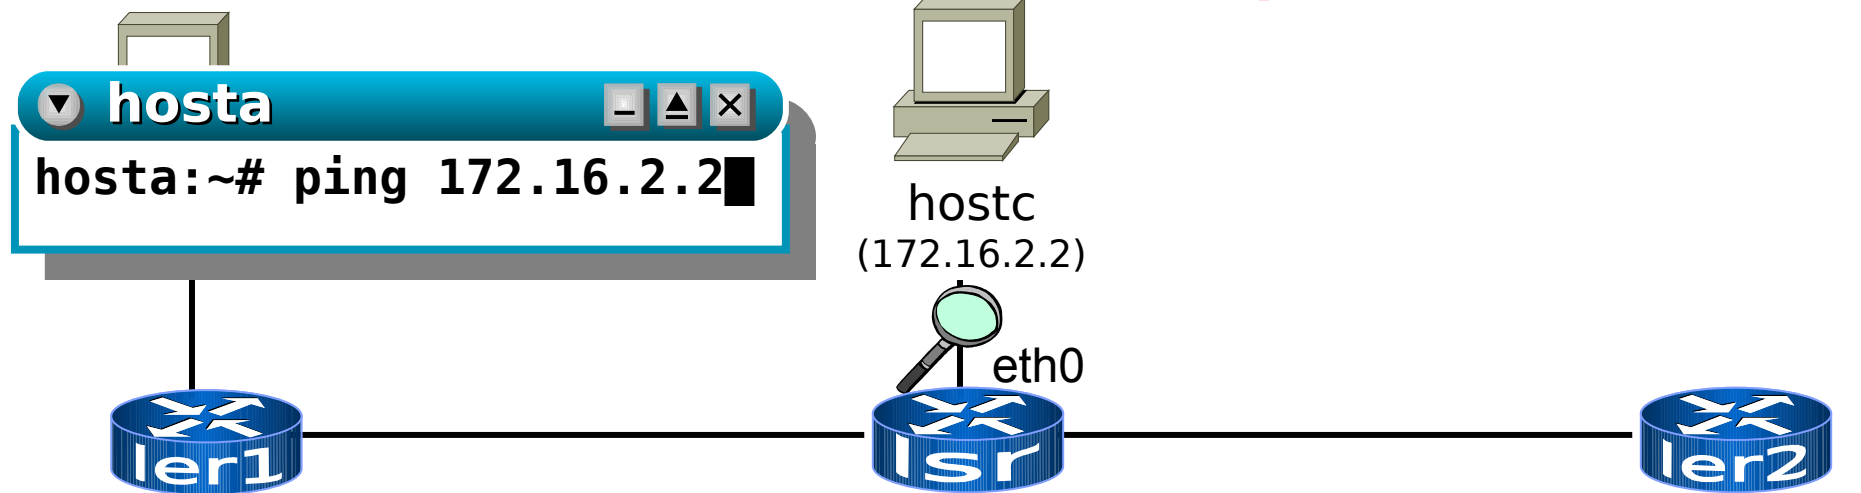
(back to)
plain ICMP
packets

Frame 1 (98 bytes on wire, 98 bytes captured)
Ethernet II, Src: ae:48:62:f5:63:c6 (ae:48:62:f5:63:c6), Dst: de:e5:92:ad:6f:e4 (de:e5:92:ad:6f:e4)
Internet Protocol, Src: 10.0.0.1 (10.0.0.1), Dst: 10.0.0.2 (10.0.0.2)
Internet Control Message Protocol

IP Traffic analysis



IP Traffic analysis



No. .	Time	Source	Destination	Protocol	Info
1	0.000000	192.168.0.1	172.16.2.2	ICMP	Echo (ping) request
2	0.000074	172.16.2.2	192.168.0.1	ICMP	Echo (ping) reply
3	0.998815	192.168.0.1	172.16.2.2	ICMP	Echo (ping) request
4	0.998933	172.16.2.2	192.168.0.1	ICMP	Echo (ping) reply
5	1.997949	192.168.0.1	172.16.2.2	ICMP	Echo (ping) request
6	1.008087	172.16.2.2	192.168.0.1	ICMP	Echo (ping) reply

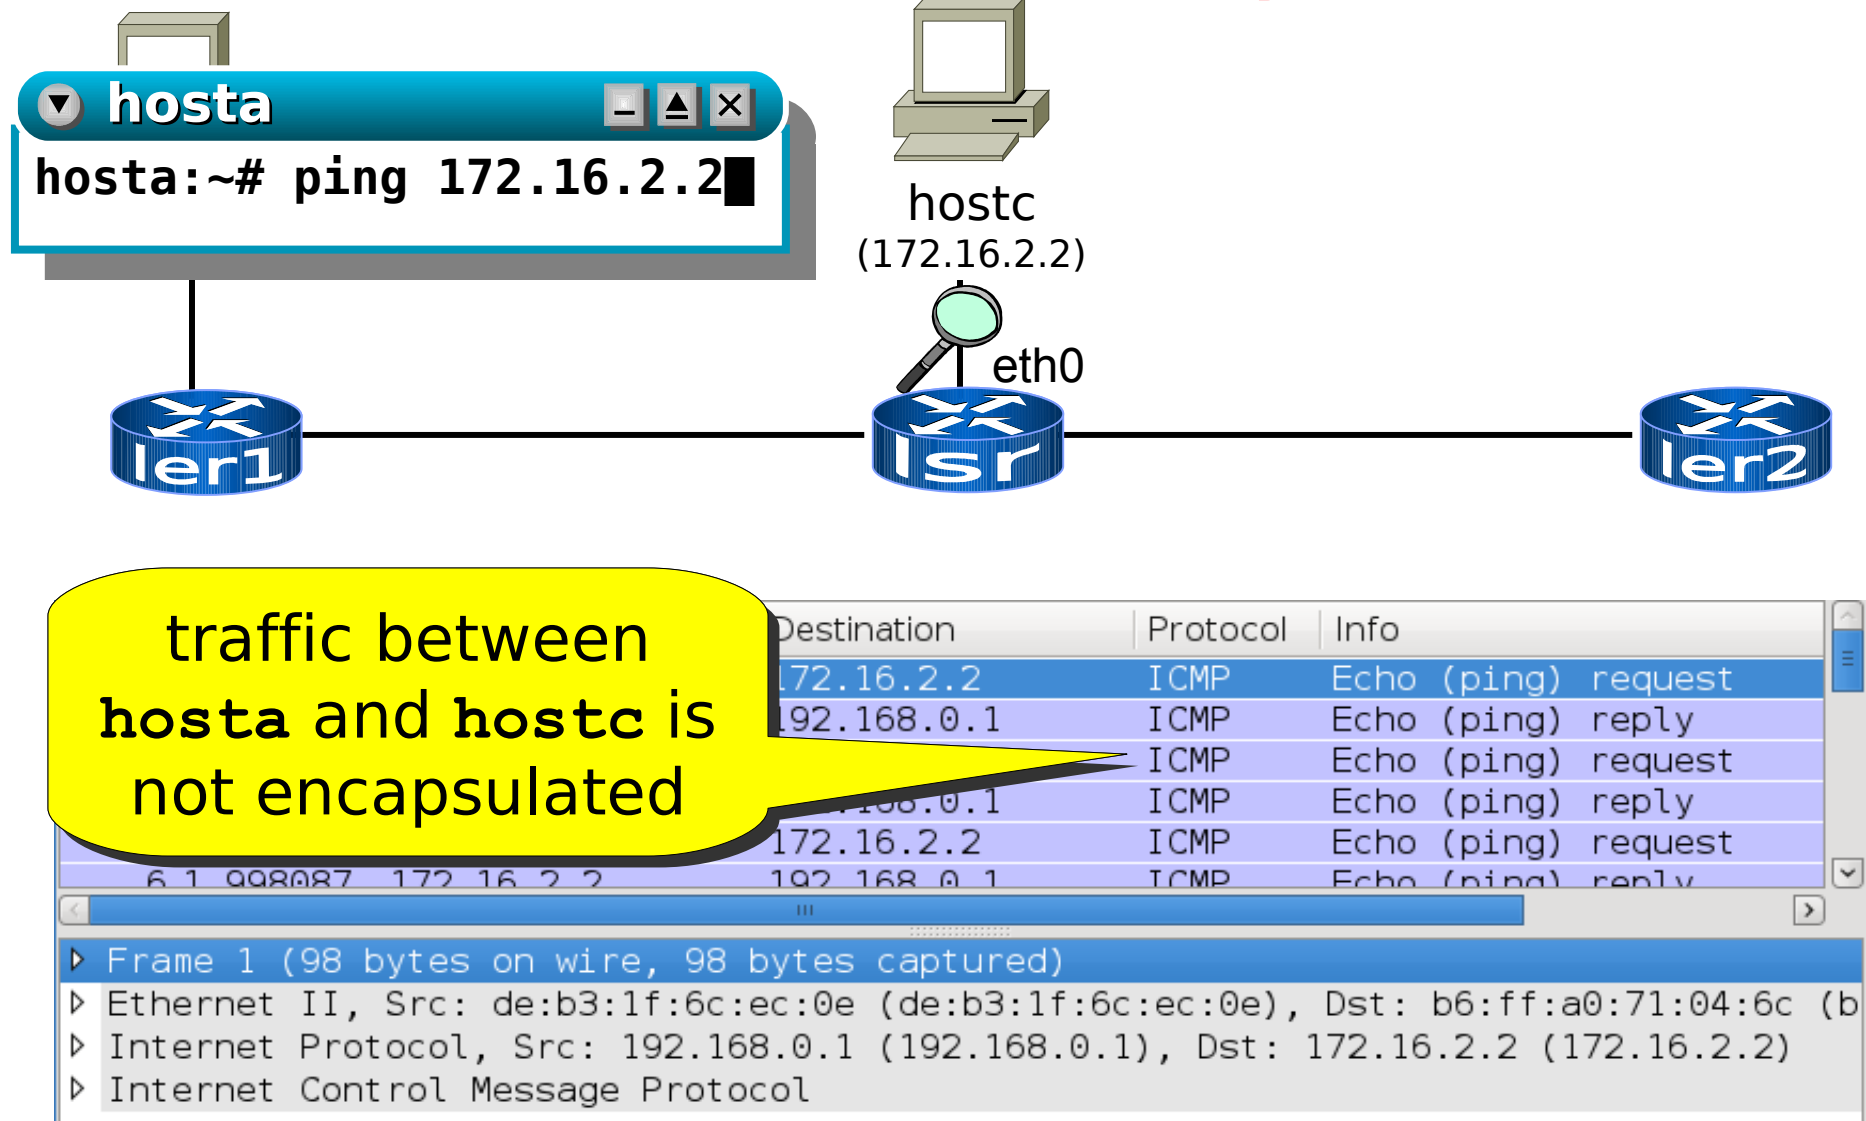
▶ Frame 1 (98 bytes on wire, 98 bytes captured)

▶ Ethernet II, Src: de:b3:1f:6c:ec:0e (de:b3:1f:6c:ec:0e), Dst: b6:ff:a0:71:04:6c (b6:ff:a0:71:04:6c)

▶ Internet Protocol, Src: 192.168.0.1 (192.168.0.1), Dst: 172.16.2.2 (172.16.2.2)

▶ Internet Control Message Protocol

IP Traffic analysis



Proposed exercises

- Sniff on all the interfaces with wireshark
- On a virtual router, try to give the commands stated in the .startup “by hand”, one by one directly on the virtual machine in order to observe the returned keys
- Try to change NHLFE tables on the running lab, adding or deleting entries (use “**mpls -help**” inside virtual machines to see the help)